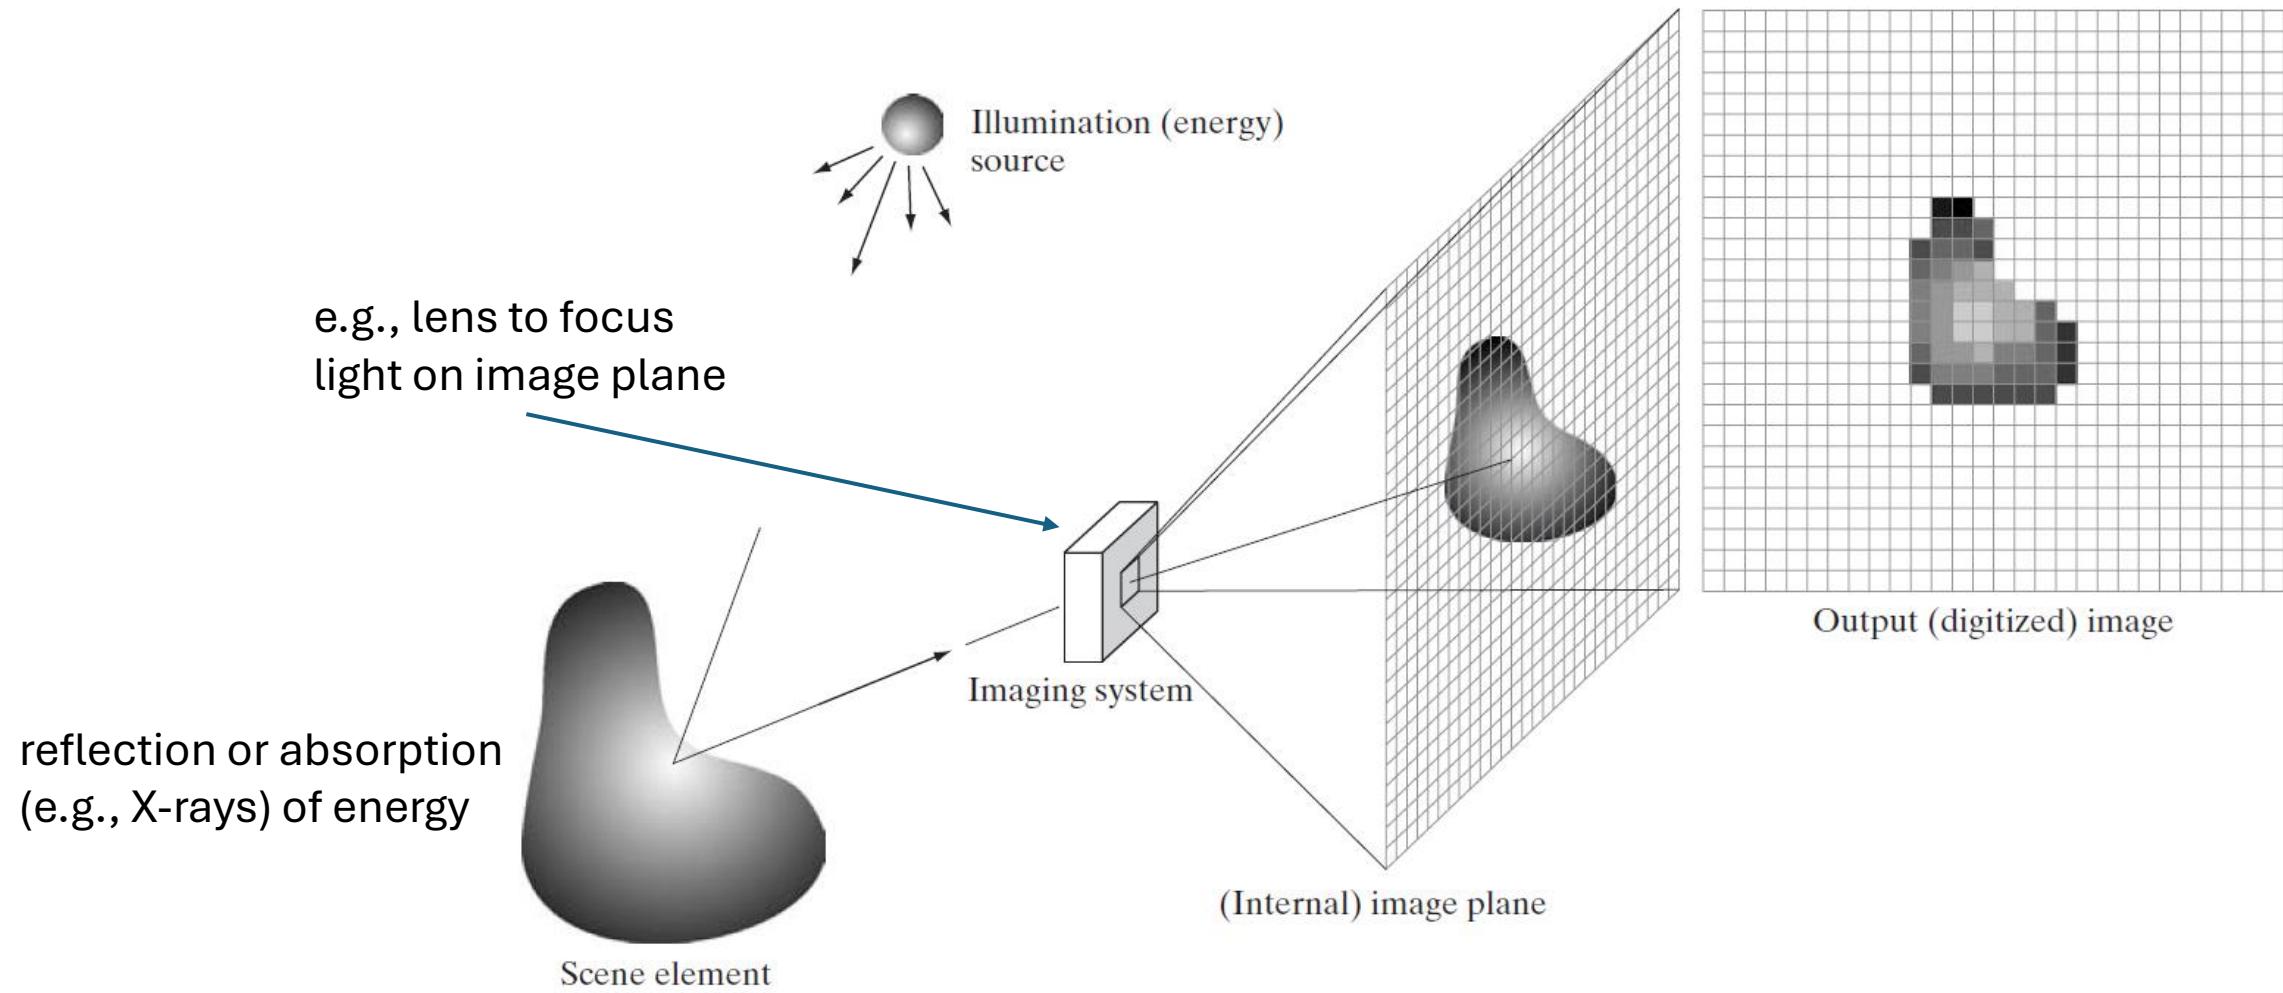


# Digital Image Processing

Deep Learning and Image Processing

# Digital Image Acquisition Process



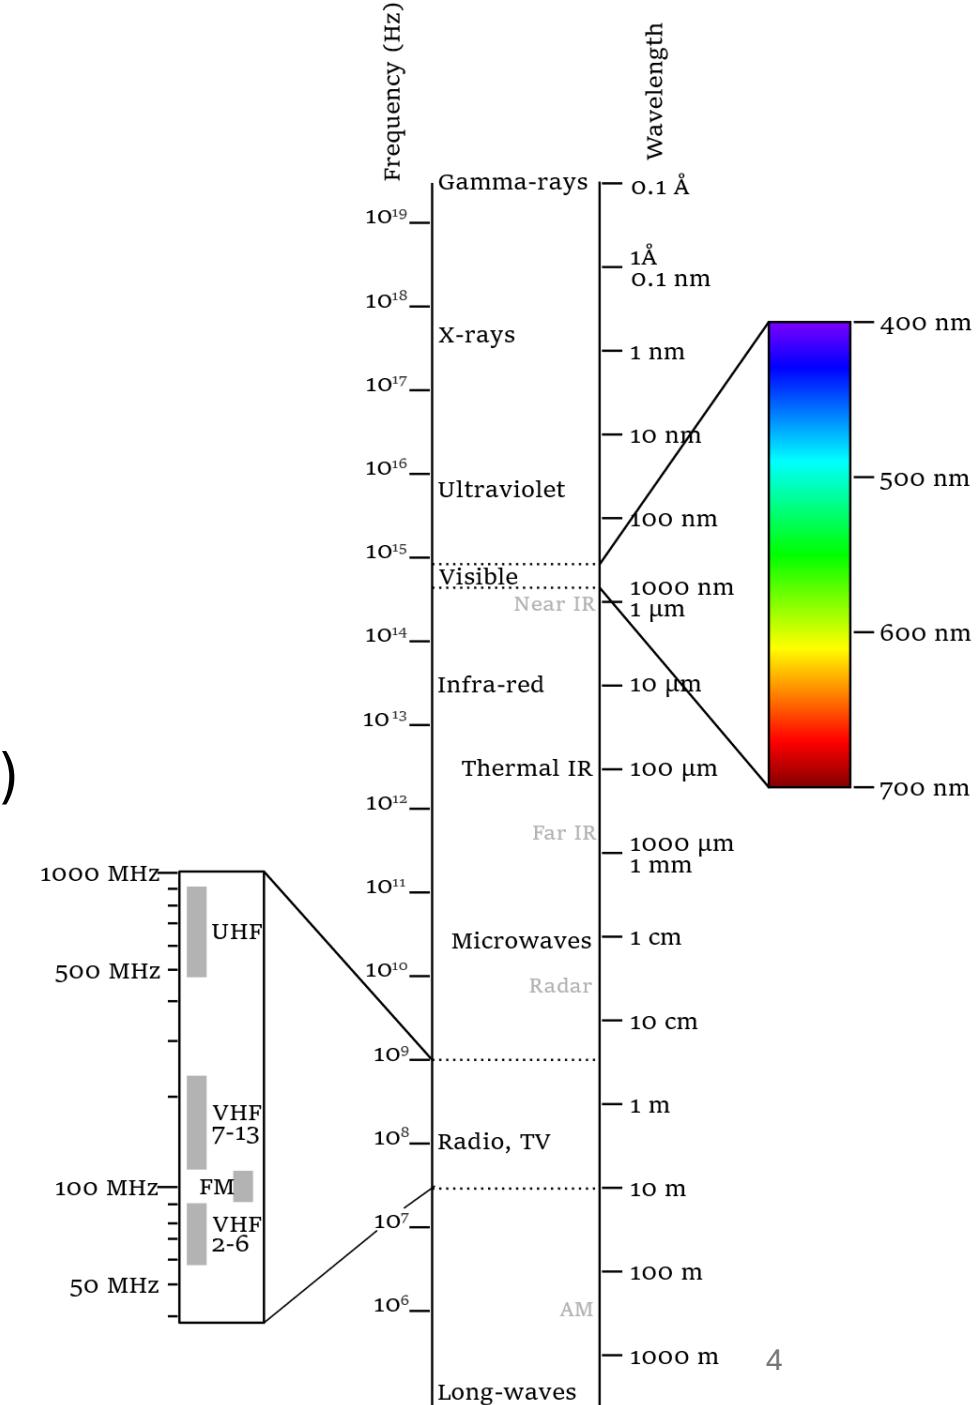
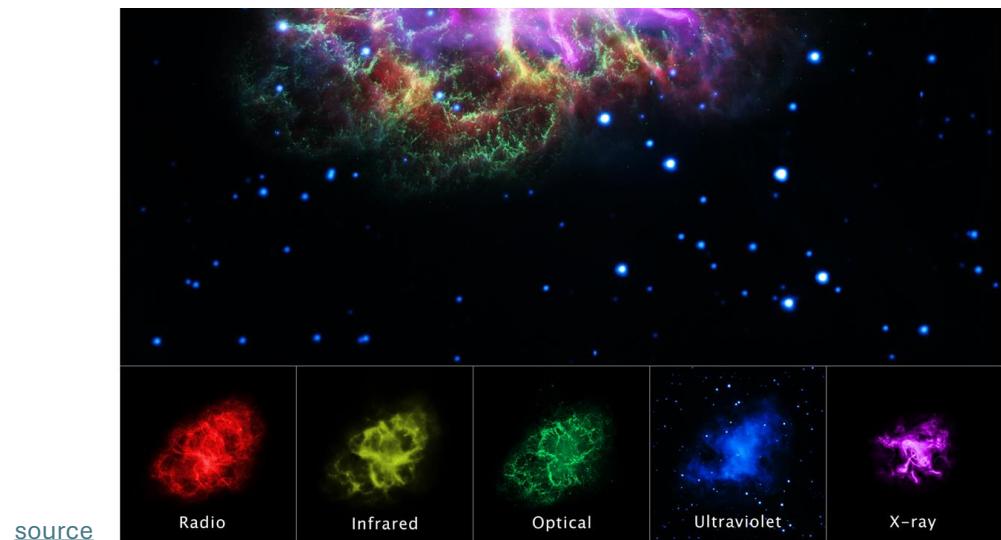
# Energy Sources

# Electromagnetic Spectrum

examples:

- visible band: digital camera
- microwaves: radar
- X-rays: computed tomography (CT)
- radio band: magnetic resonance imaging (MRI)
- gamma rays: positron emission tomography (PET)

Crab Nebula  
in multiple  
wavelengths



# Lidar vs Cameras for Self-Driving Cars

Lidar:

- similar to radar, but using lasers
- very accurate



camera:

- human-like perception
- cost-effective

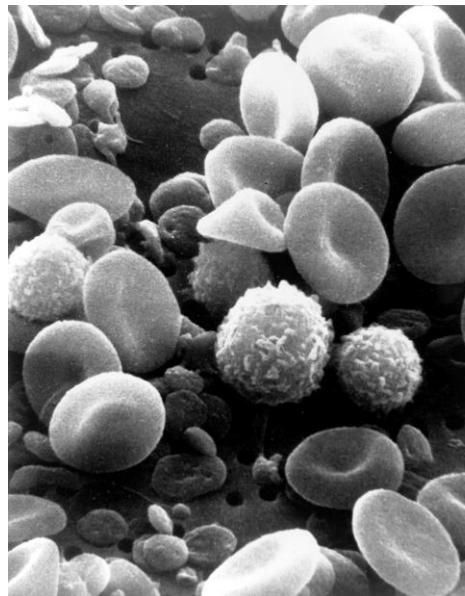


# Examples for other Energy Sources

ultrasound

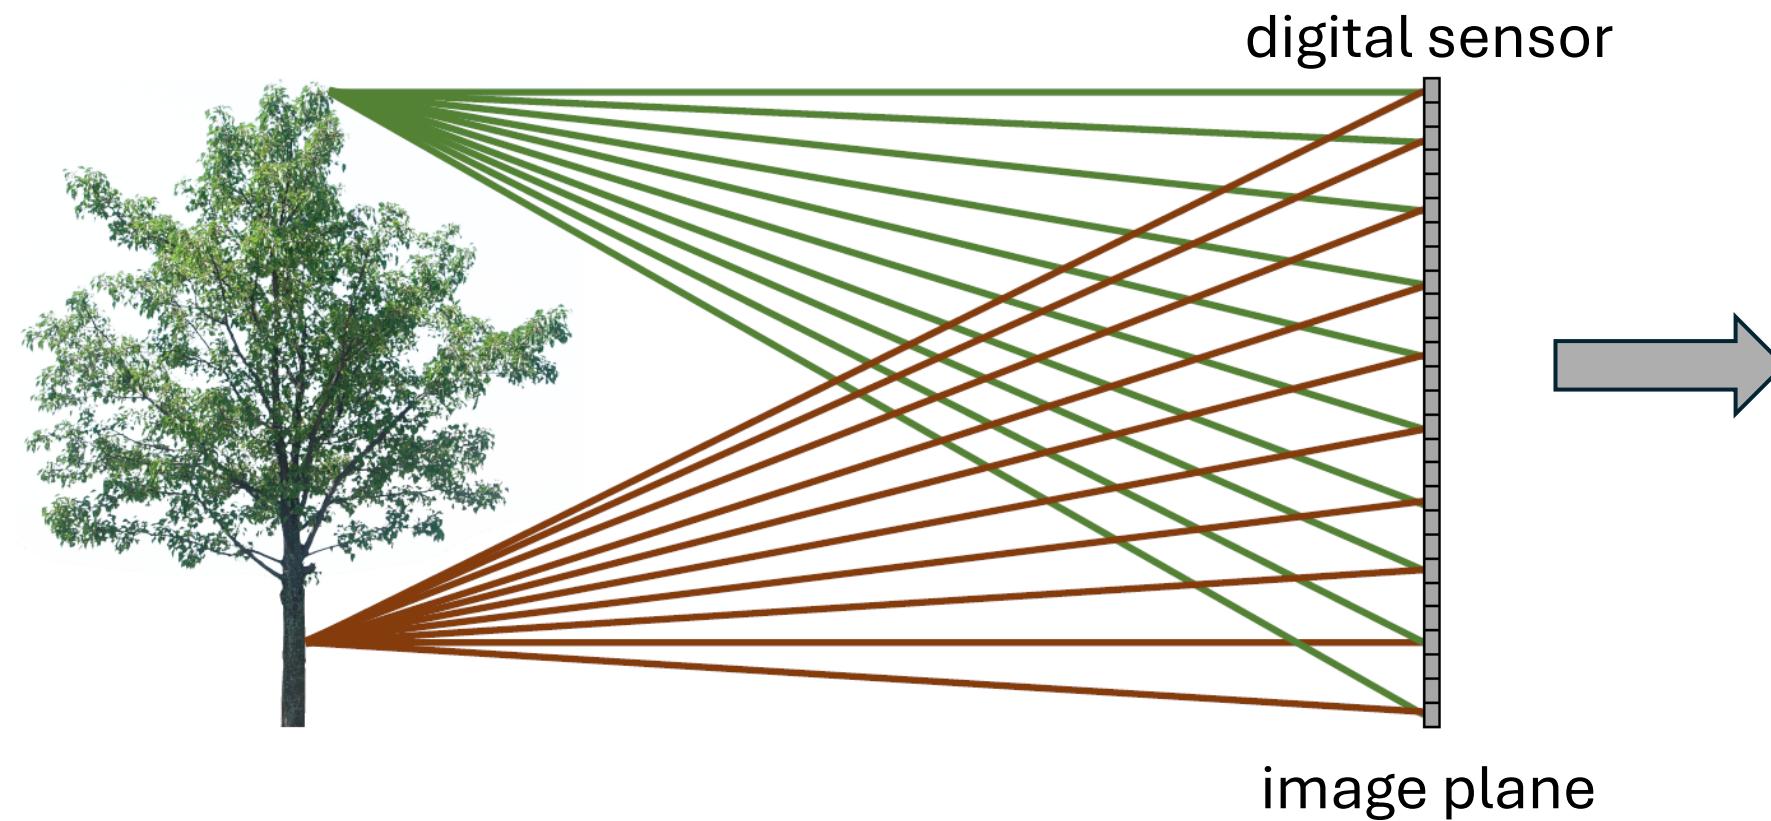


electron beams  
(electron microscopy)

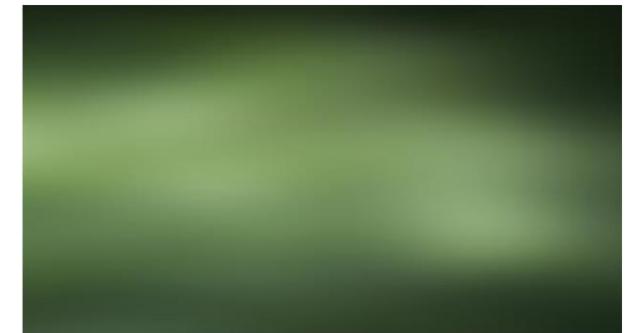


# Image Formation

# Bare-Sensor Imaging

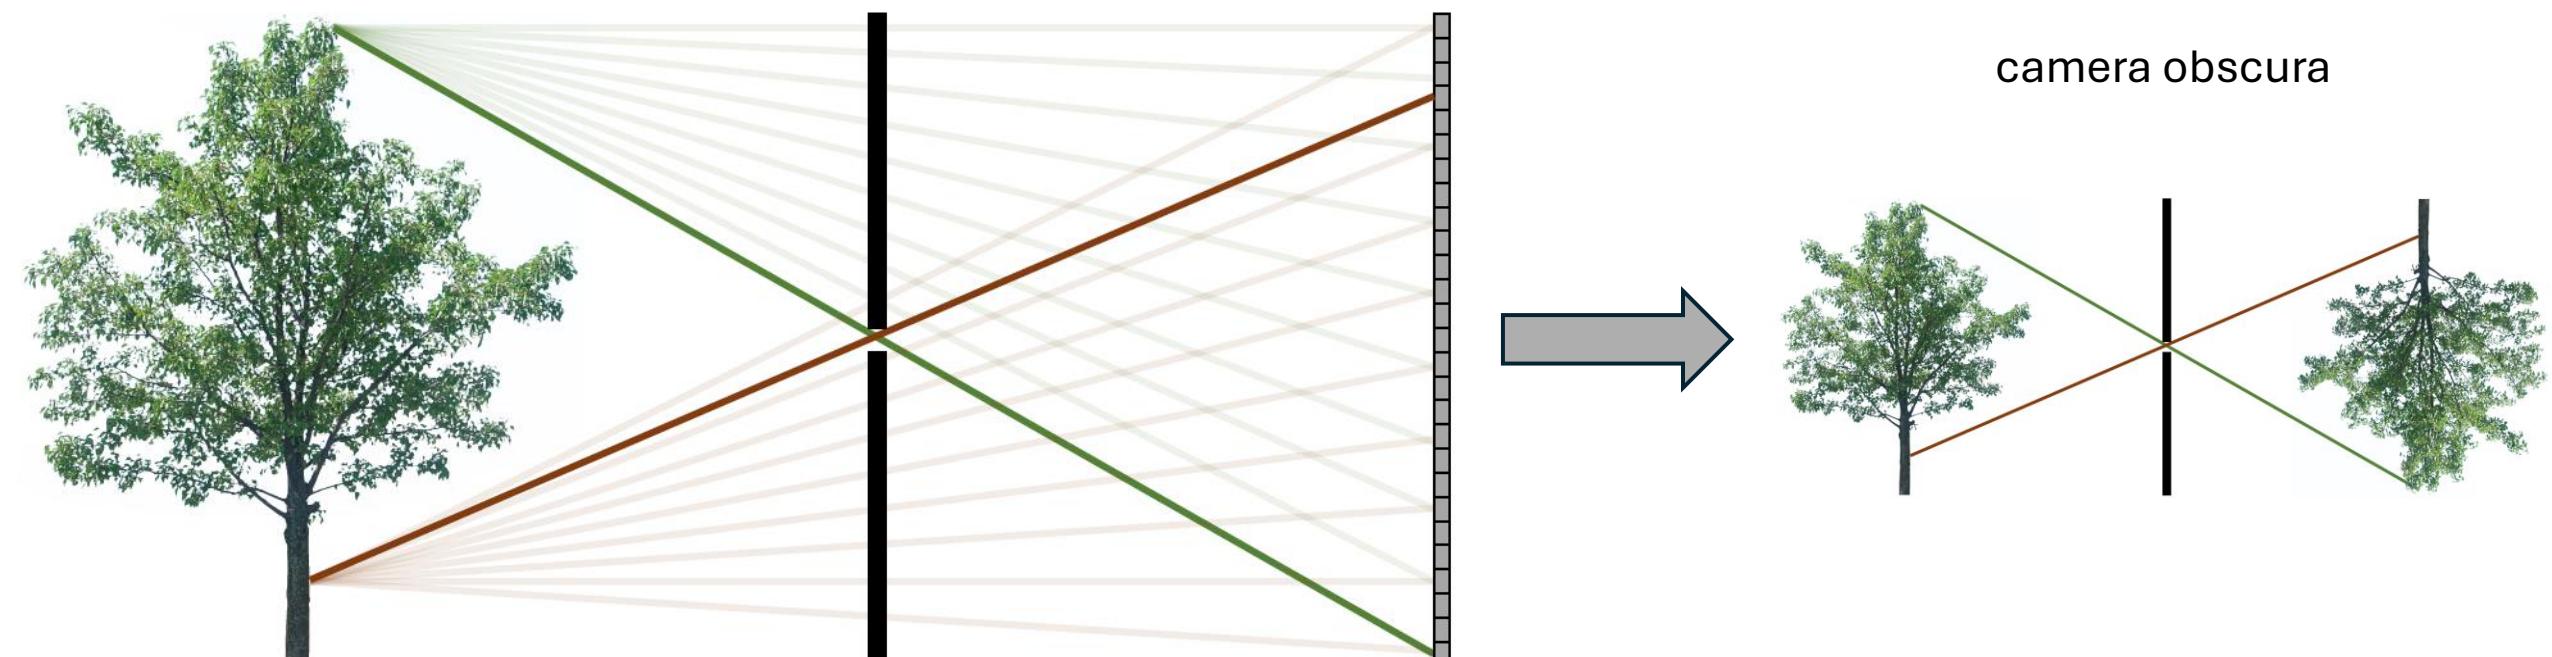


all scene points contribute  
to all sensor pixels:



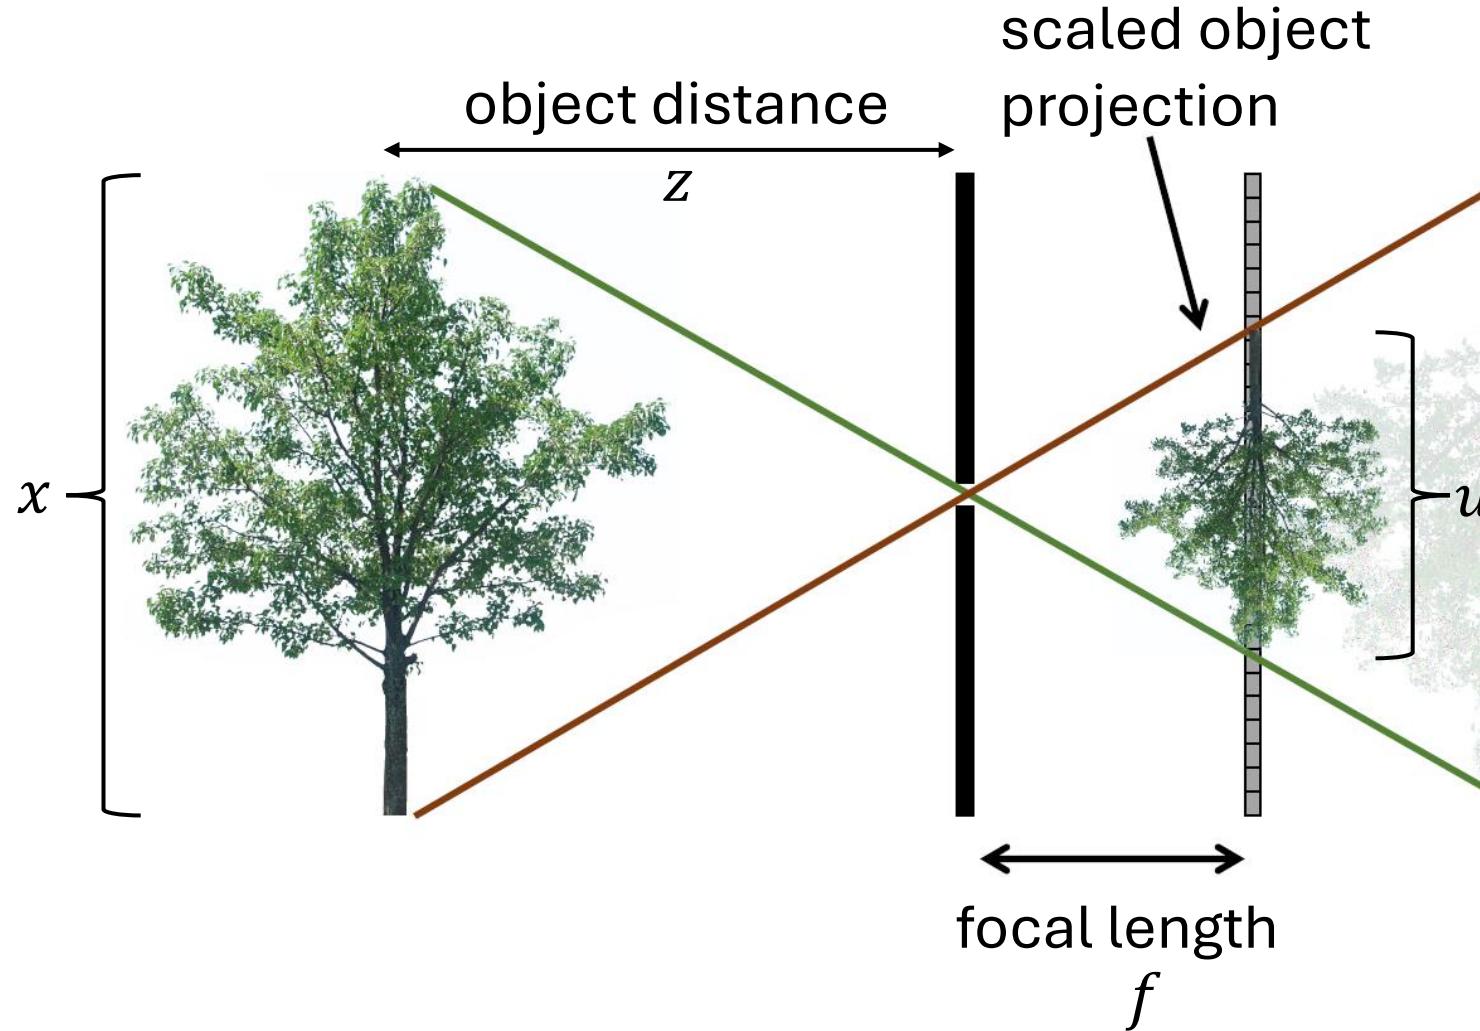
# Pinhole Camera

camera: barrier with  
pinhole (aperture)



each scene point contributes  
to one pixel only

# Pinhole Camera Model



3D world coordinates  $(x, y, z)$  to  
2D image coordinates  $(u, v)$ :

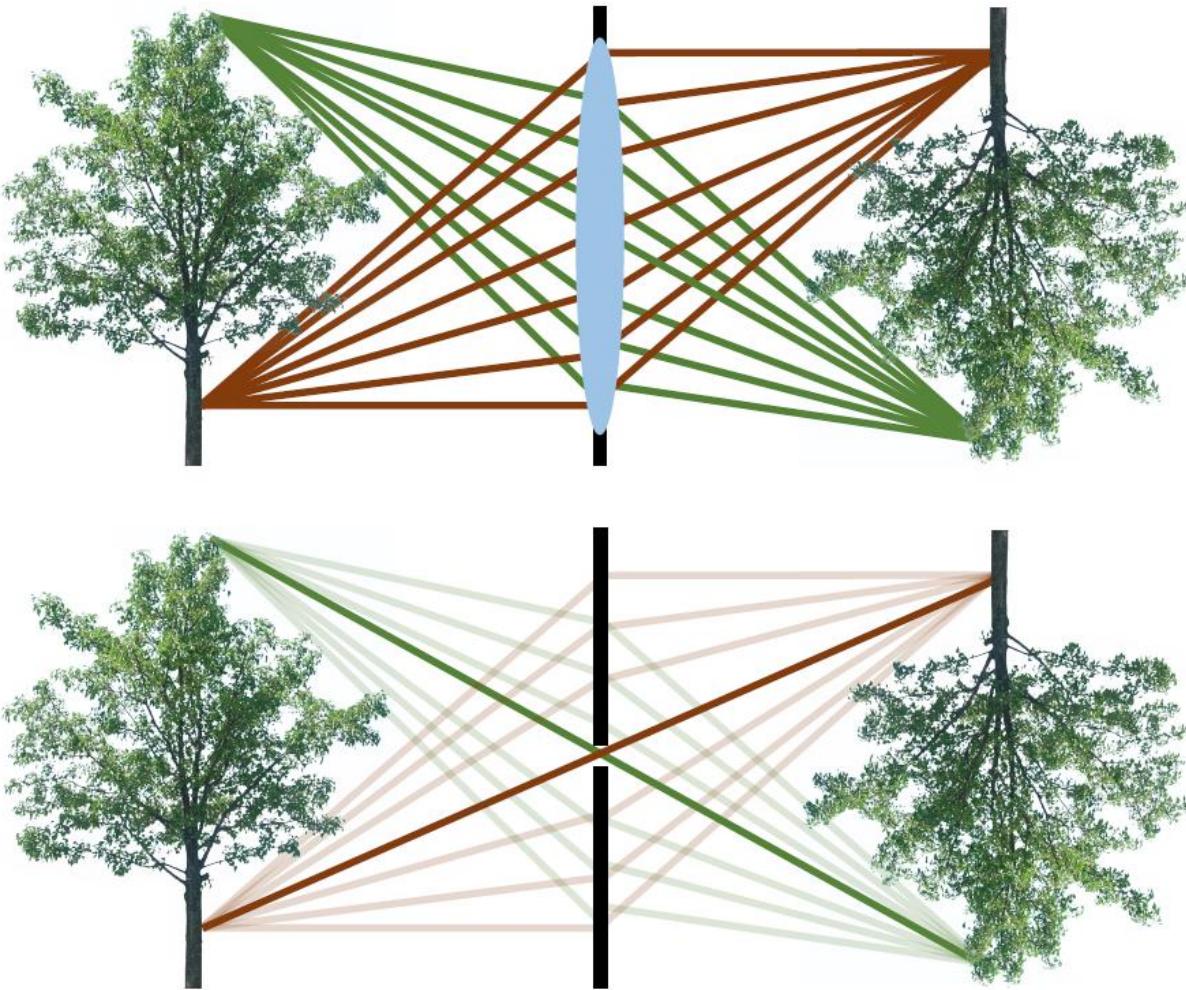
$$\begin{pmatrix} u \\ v \end{pmatrix} = -\frac{f}{z} \begin{pmatrix} x \\ y \end{pmatrix}$$

losing depth



(e.g., stereo vision to get it back)

# Lens Camera

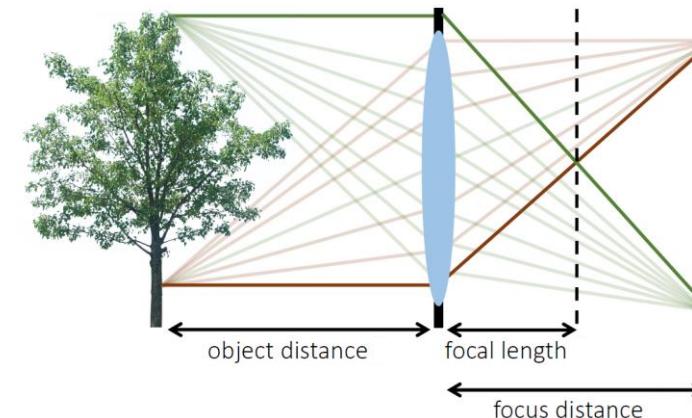


ideal pinhole is infinitesimally small: larger diameter/aperture makes image blurrier (but also allows more light to pass through: higher intensity of image pixels)

→ use optical lens

lens camera can be described with pinhole camera model, if:

- only central rays are considered
- lens camera is in focus (focus distance of lens camera equal to focal length of pinhole camera)



# Digital Image Sensors

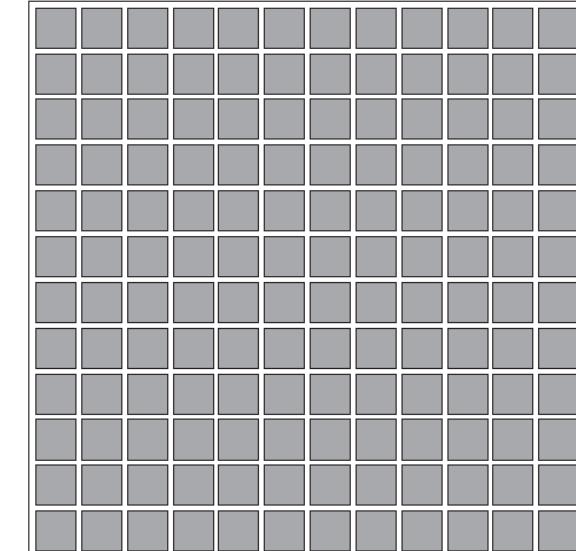
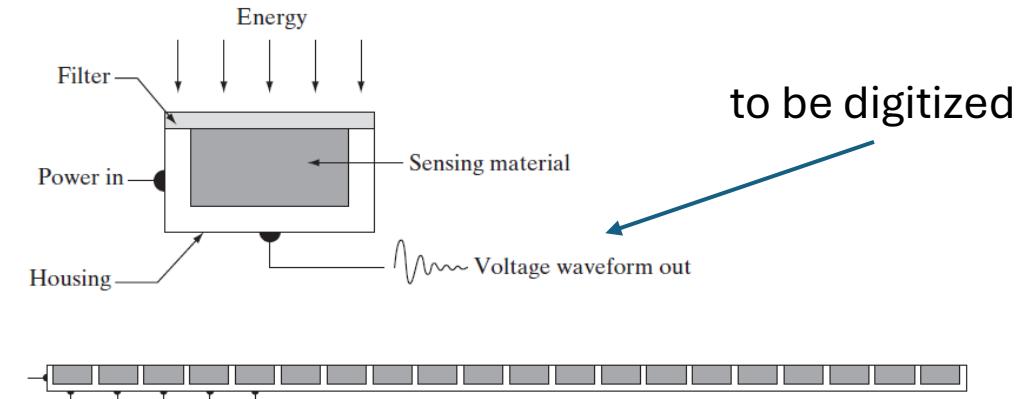
# Photons to Electrons

largely replaced chemical and analog imaging

single sensor: photodiode

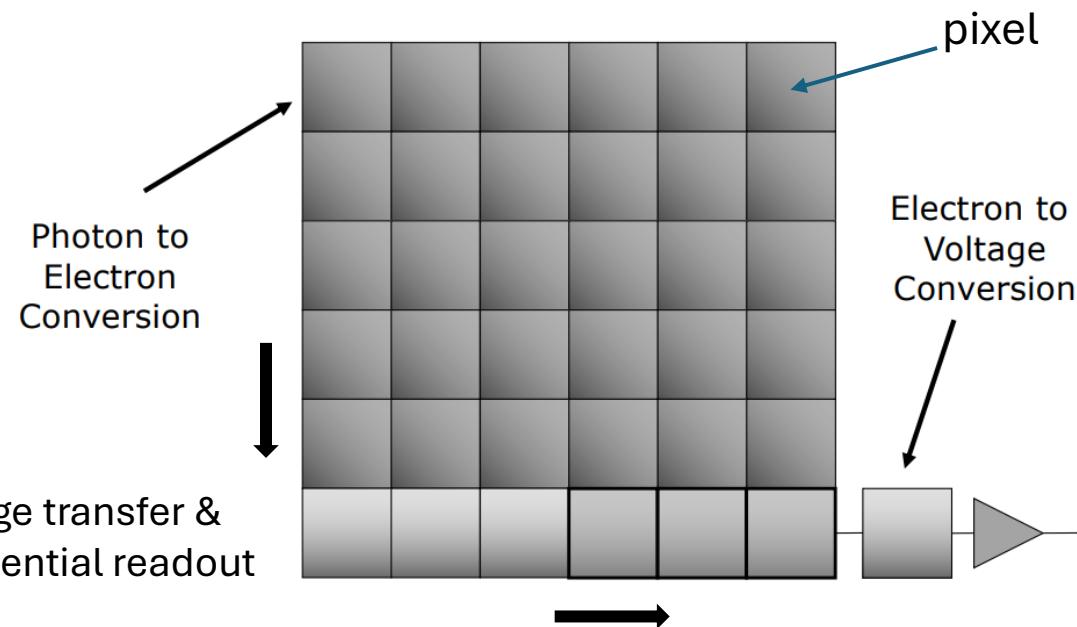
sensor strips: motion perpendicular to strip for imaging in other direction, used in CT

sensor arrays: used in digital cameras (in focal plane, outputs proportional to integral of light received at each sensor)



# Most Popular Sensors: CCD and CMOS

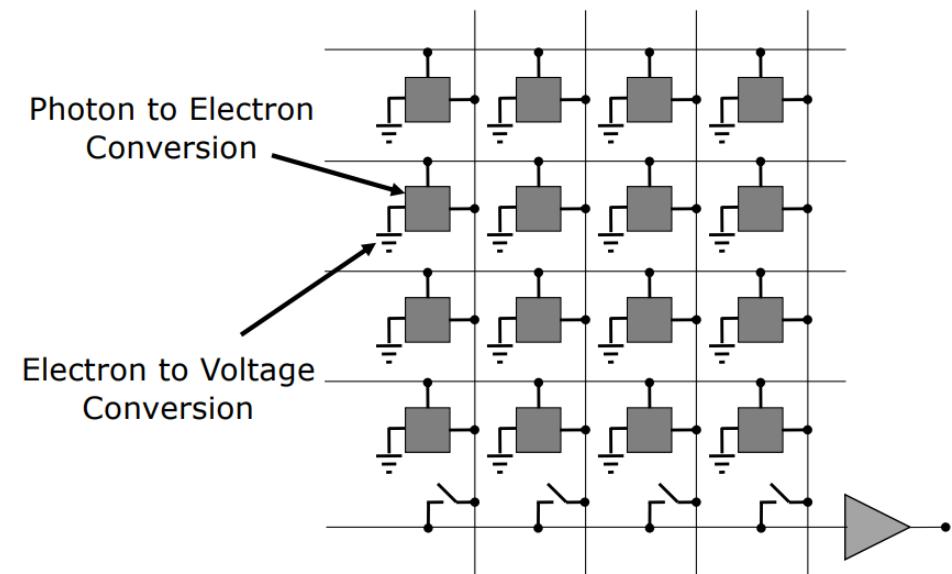
CCD: Charge Coupled Device



larger light-sensitive areas, lower noise level  
→ used in high-end video cameras

active pixel sensor

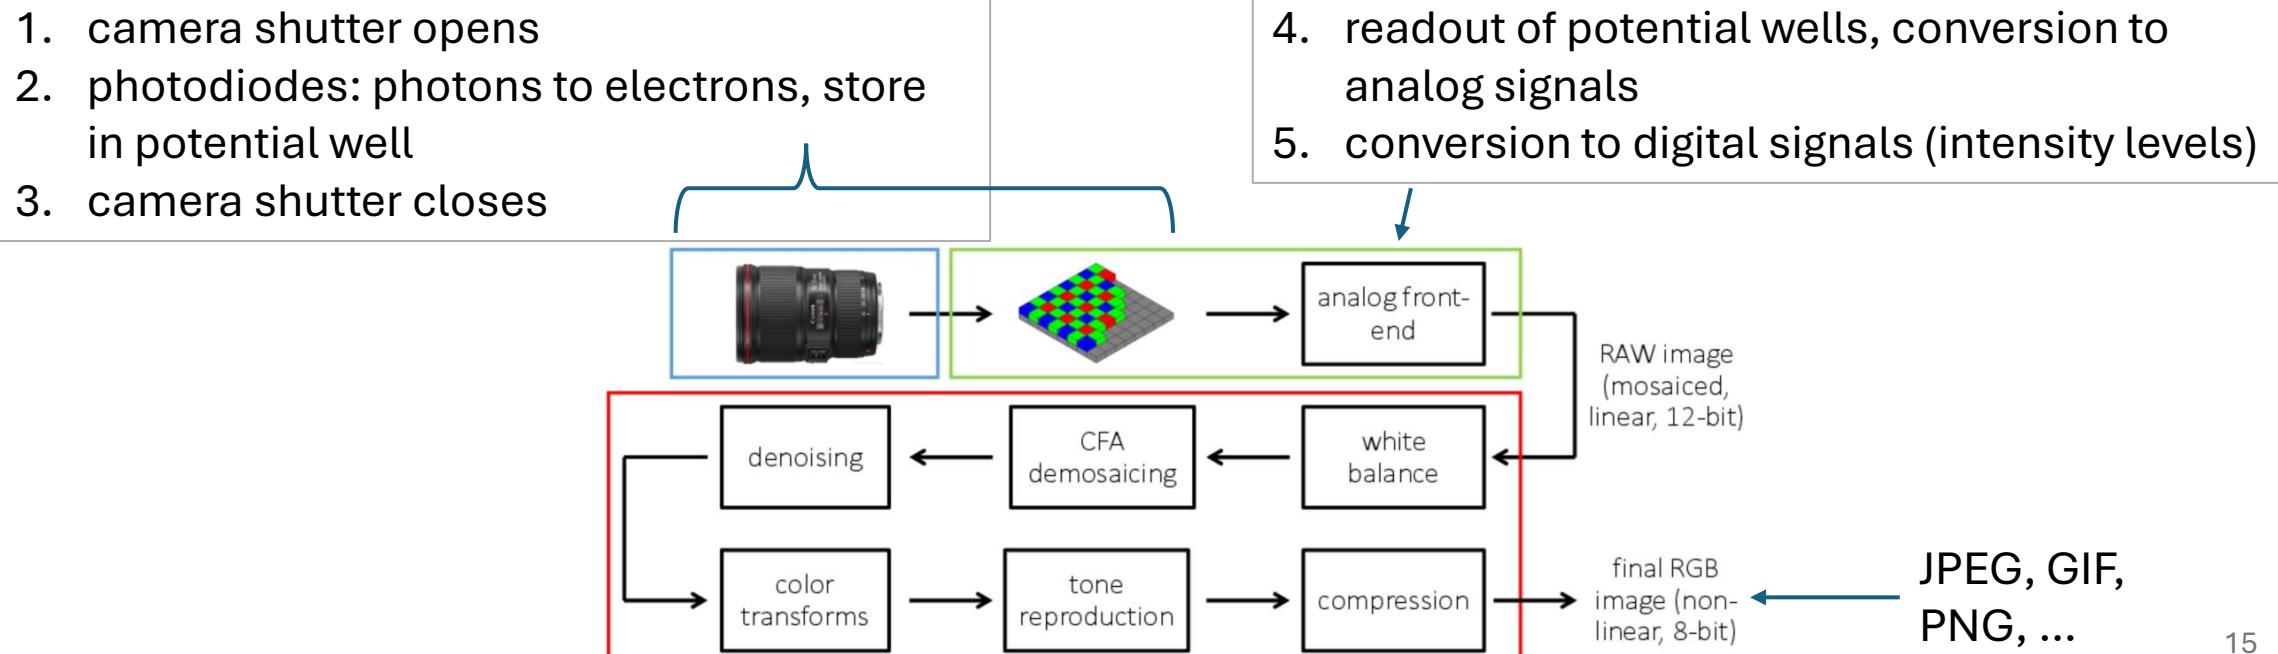
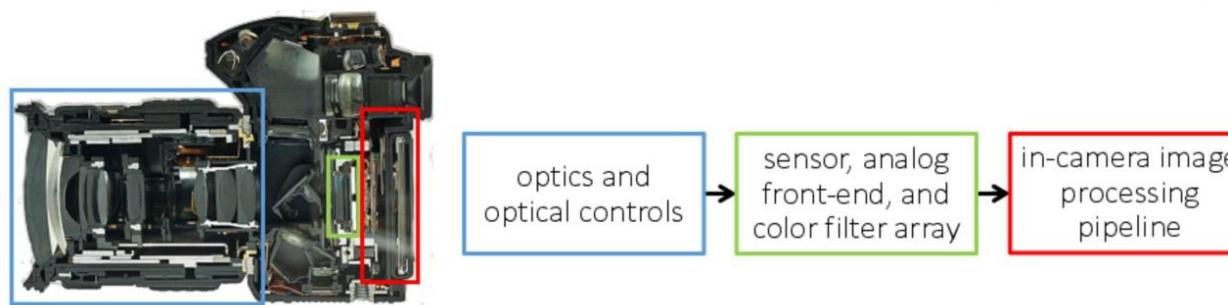
CMOS: Complementary Metal-Oxide Semiconductor



lower power consumption, typically smaller  
→ used in most consumer product cameras

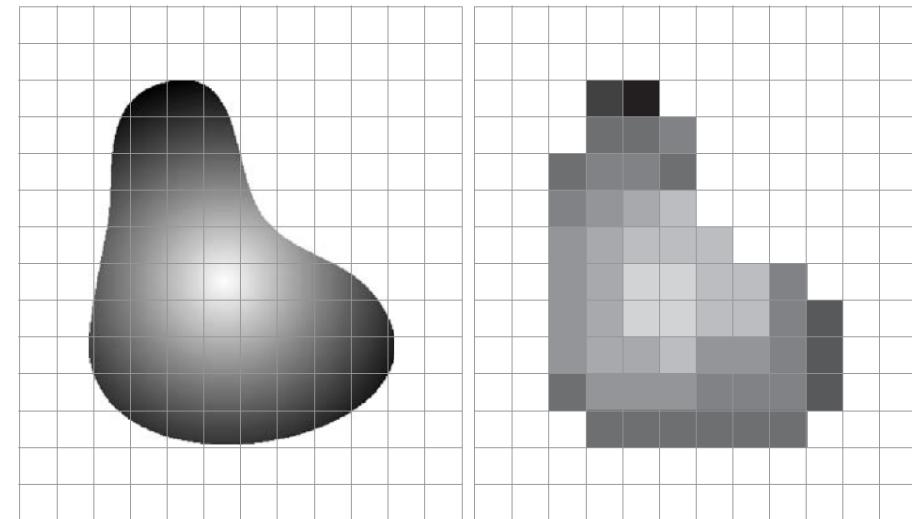
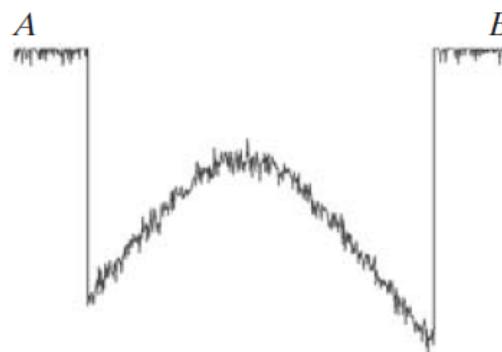
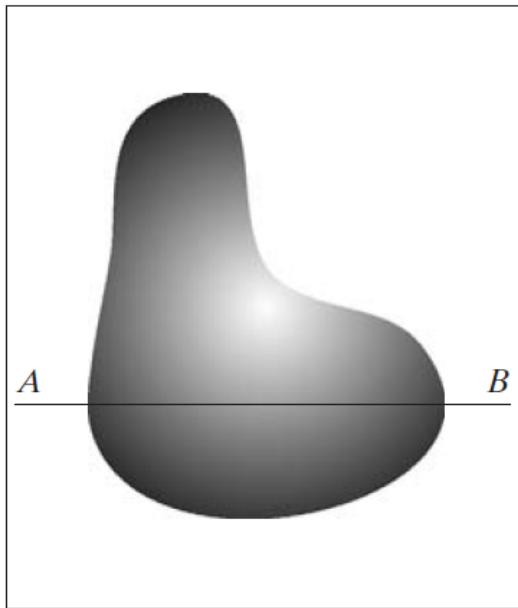
no individual silicon cells, but microlenses on top of each pixel to focus light on sensitive areas

# The modern photography pipeline

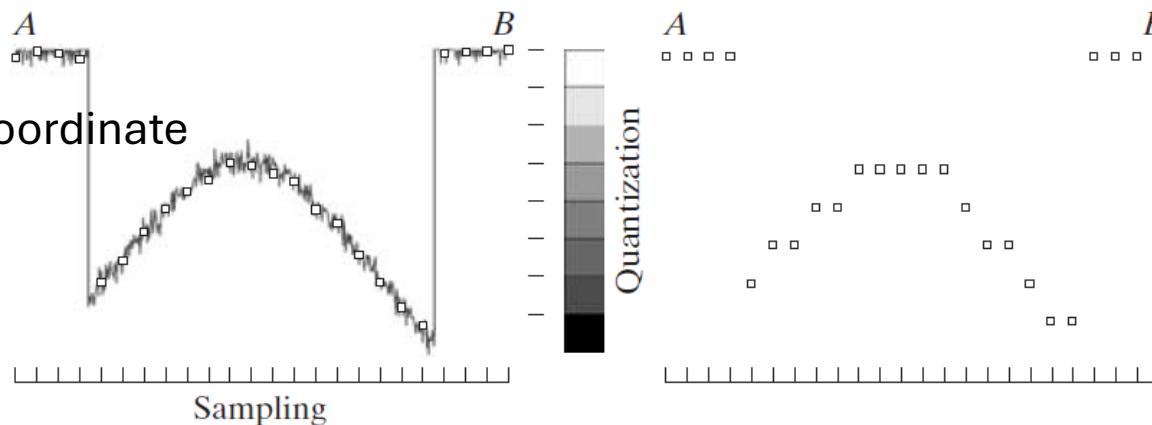


# Pixels & Intensity

# Digitizing: Sampling & Quantization



digitizing the coordinate  
values: pixels



digitizing the amplitude values:  
discrete intensity levels

# Discrete Intensity Levels

image data: 2D grid (matrix) of pixels with different intensity values

intensity resolution:

common to use one byte ( $2^8$ ) to express possible intensity values (8-bit image)

→ 0 = black, 255 = white

storage requirement for, e.g.,  $1024 \times 1024$  8-bit image (grayscale): ~1 MB  
( $\times 3$  for RGB color images)

0	2	15	0	0	11	10	0	0	0	9	9	0	0	0	0
0	0	0	4	60	157	236	255	255	177	95	61	32	0	0	29
0	10	16	119	238	255	244	245	243	250	249	255	222	103	10	0
0	14	170	255	255	244	254	255	253	245	255	249	253	251	124	1
2	98	255	228	255	251	254	211	141	116	122	215	251	238	255	49
13	217	243	255	155	33	226	52	2	0	10	13	232	255	255	36
16	229	252	254	49	12	0	0	7	7	0	70	237	252	235	62
6	141	245	255	212	25	11	9	3	0	115	236	243	255	137	0
0	87	252	250	248	215	60	0	1	121	252	255	248	144	6	0
0	13	113	255	255	245	255	182	181	248	252	242	208	36	0	19
1	0	5	117	251	255	241	255	247	255	241	162	17	0	7	0
0	0	0	4	58	251	255	246	254	253	255	120	11	0	1	0
0	0	4	97	255	255	255	248	252	255	244	255	182	10	0	4
0	22	206	252	246	251	241	100	24	113	255	245	255	194	9	0
0	111	255	242	255	158	24	0	0	6	39	255	232	230	56	0
0	218	251	250	137	7	11	0	0	0	2	62	255	250	125	3
0	173	255	255	101	9	20	0	13	3	13	182	251	245	61	0
0	107	251	241	255	230	98	55	19	118	217	248	253	255	52	4
0	18	146	250	255	247	255	255	255	249	255	240	255	129	0	5
0	0	23	113	215	255	250	248	255	255	248	248	118	14	12	0
0	0	6	1	0	52	153	233	255	252	147	37	0	0	4	1
0	0	5	5	0	0	0	0	0	14	1	0	6	6	0	0

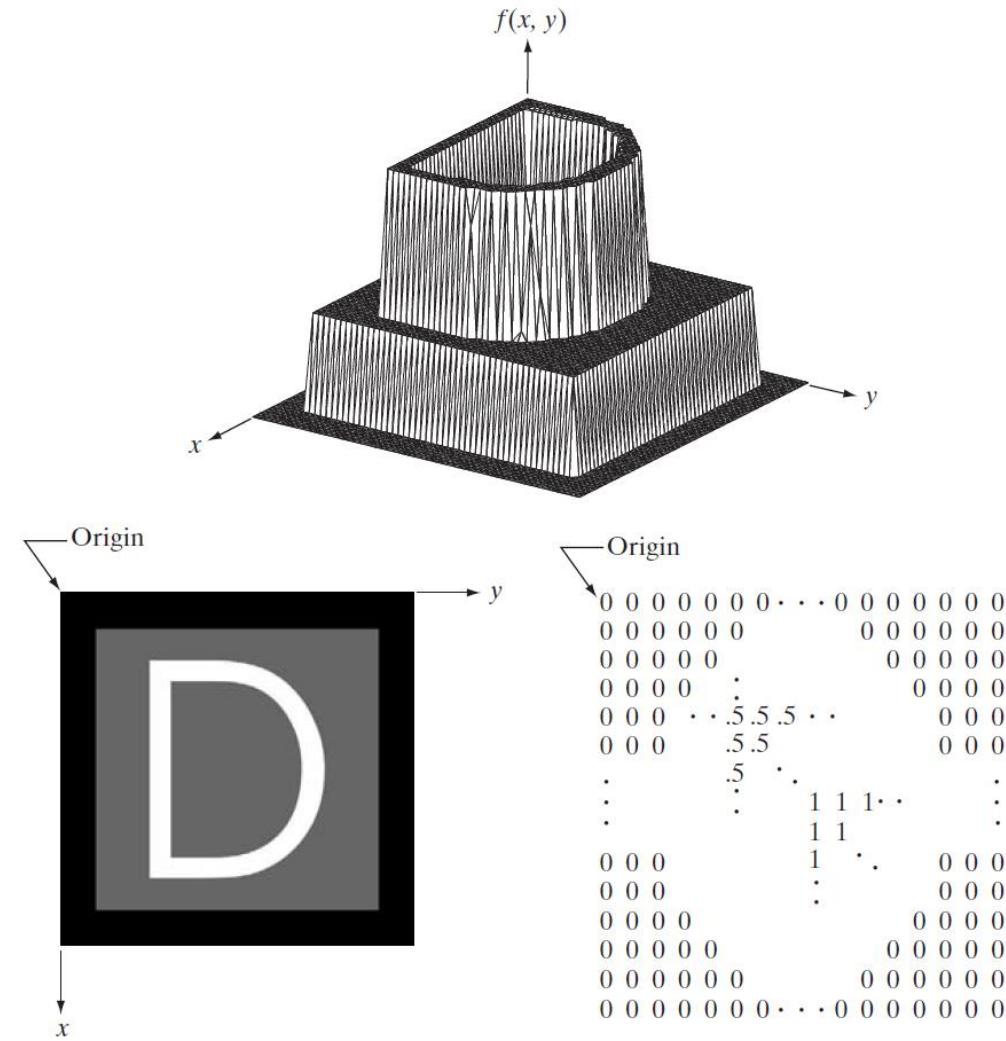
[source](#)

# Image as Two-Dimensional Function

image can be seen as function  $f(x, y)$ :  
intensity at position  $(x, y)$

digital image: discrete (sampled and quantized) version of this function

$$f(x, y) = \begin{bmatrix} f(0, 0) & f(0, 1) & \cdots & f(0, N - 1) \\ f(1, 0) & f(1, 1) & \cdots & f(1, N - 1) \\ \vdots & \vdots & & \vdots \\ f(M - 1, 0) & f(M - 1, 1) & \cdots & f(M - 1, N - 1) \end{bmatrix}$$

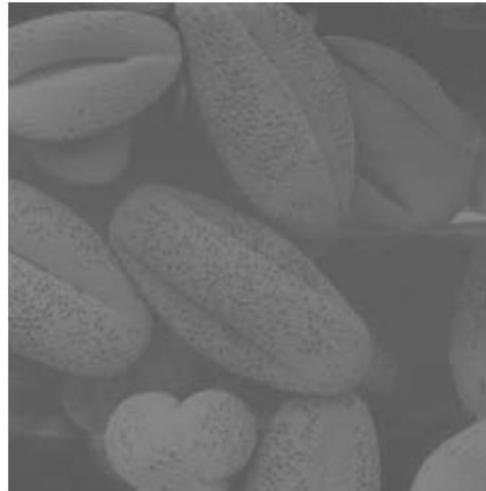


# Contrast

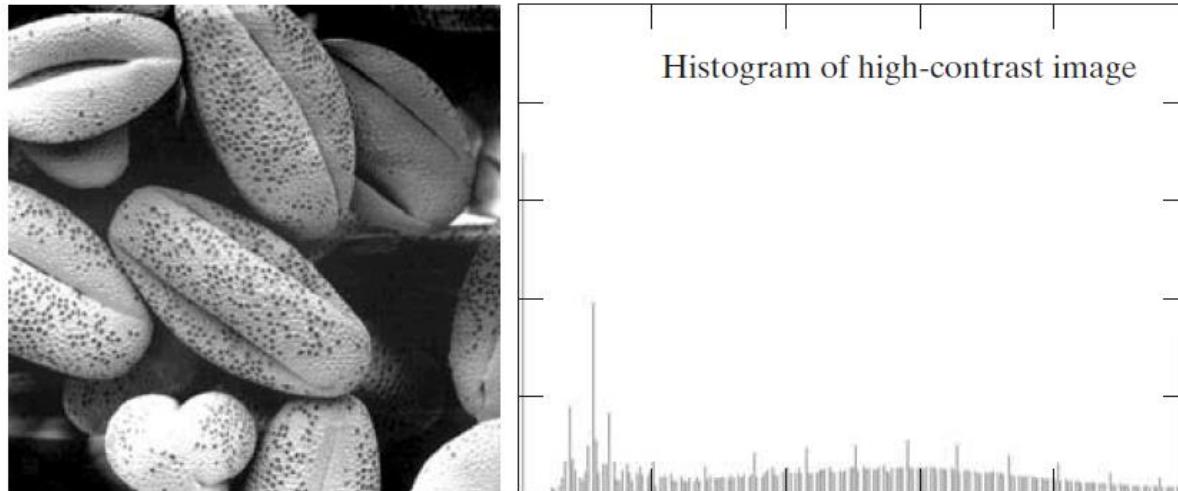
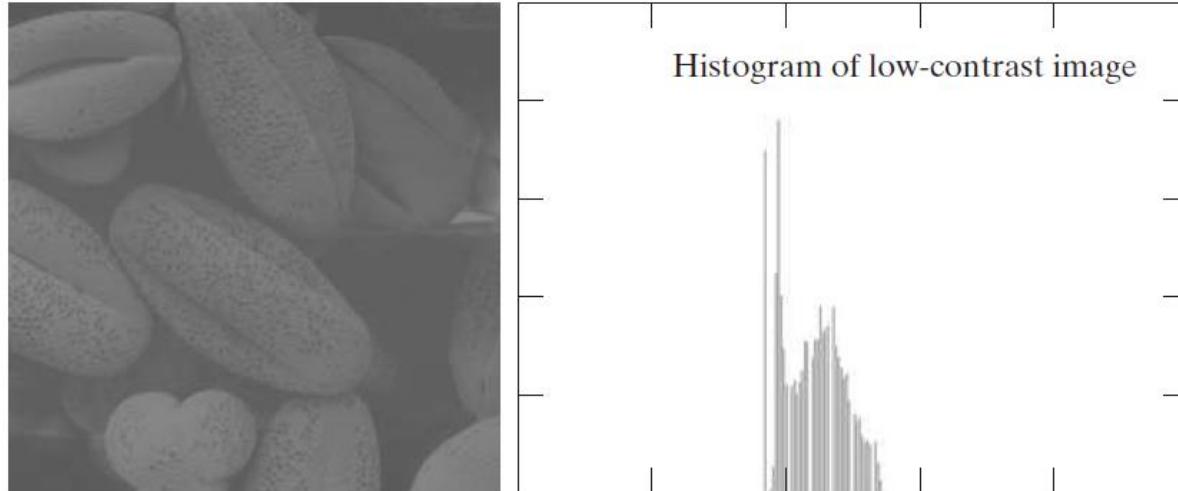
range of intensity levels:

- maximum measurable intensity determined by saturation
- minimum detectable intensity determined by noise

contrast: difference in intensity between highest and lowest intensity levels in an image



intensity histograms:

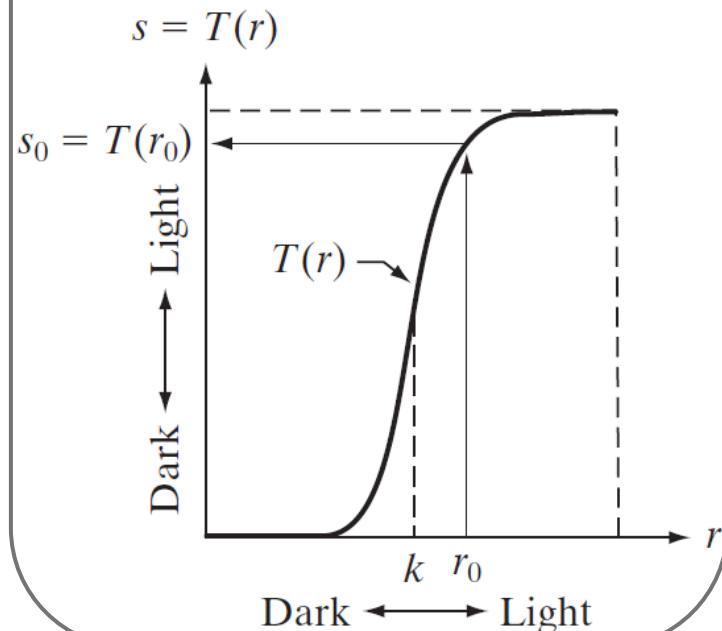


lighten or darken image simply by adding or subtracting a fixed value from the pixel intensities

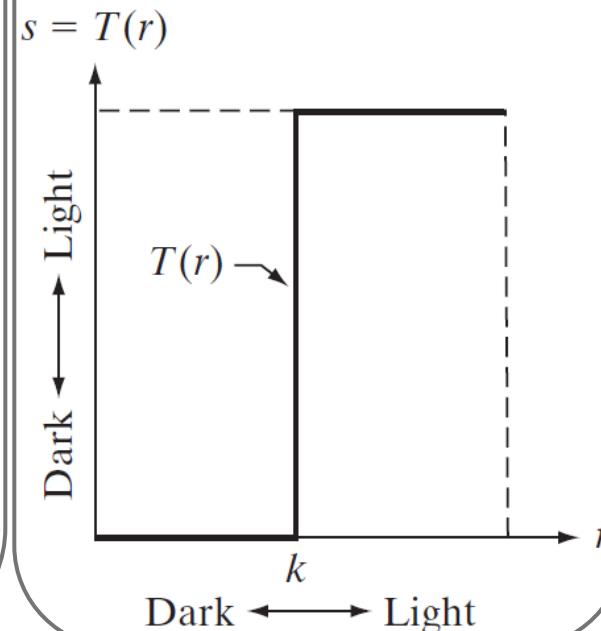
# Intensity Transformations

e.g., image enhancement through contrast manipulation

**contrast stretching**  
(expanding range of intensity levels):



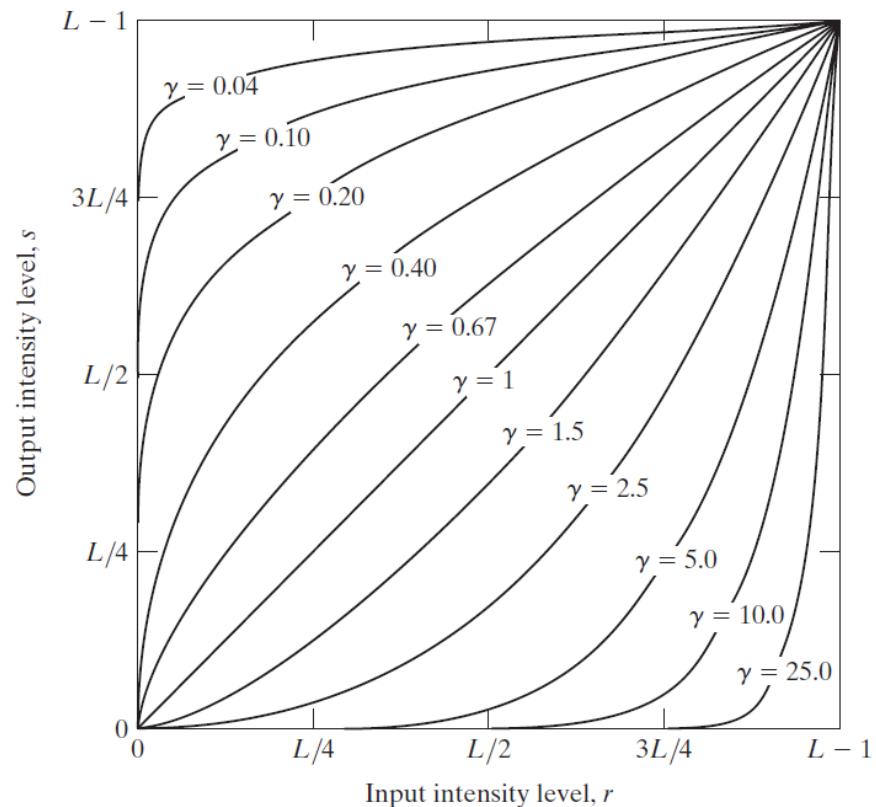
**thresholding**  
(differentiation of foreground and background regions):



# Power-Law Transformations

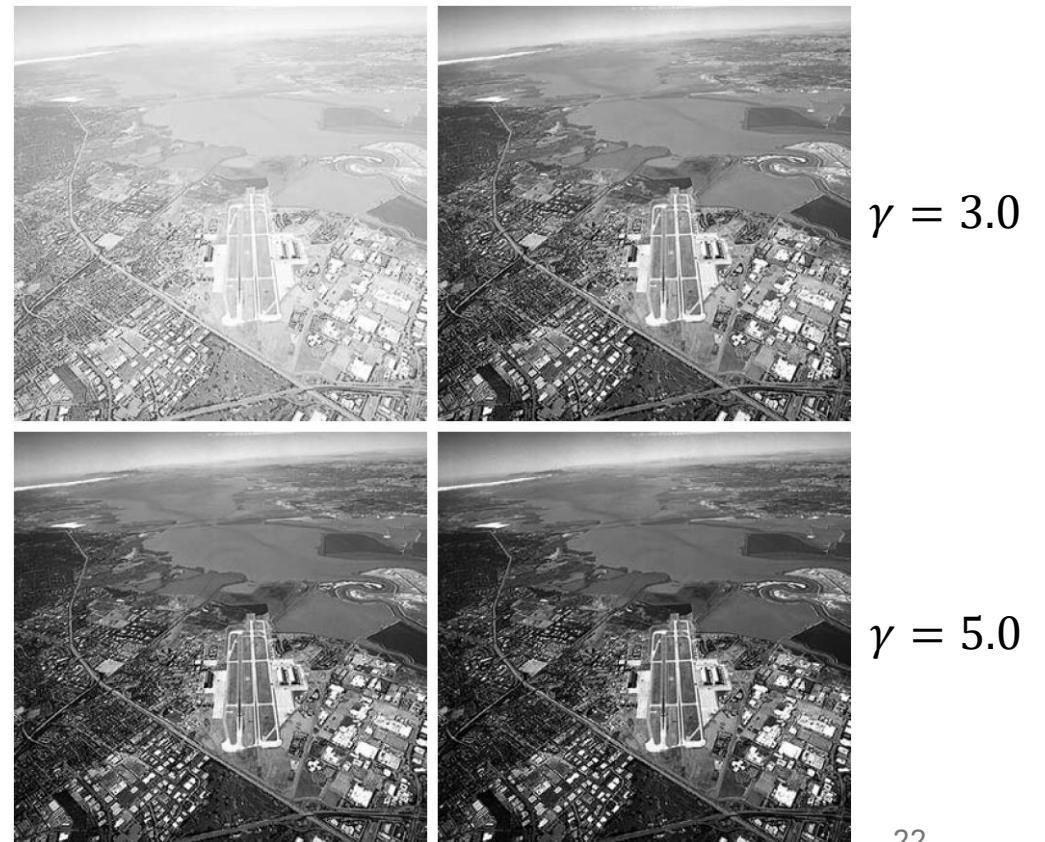
simultaneous spreading and compression of intensity levels

e.g., gamma correction:



$$s = r^\gamma$$

$$\gamma = 4.0$$

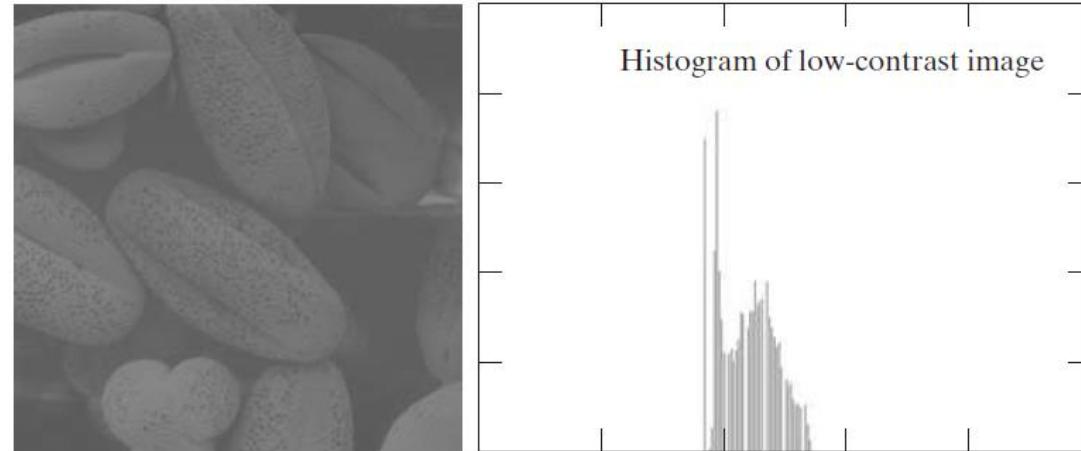


# Histogram Equalization

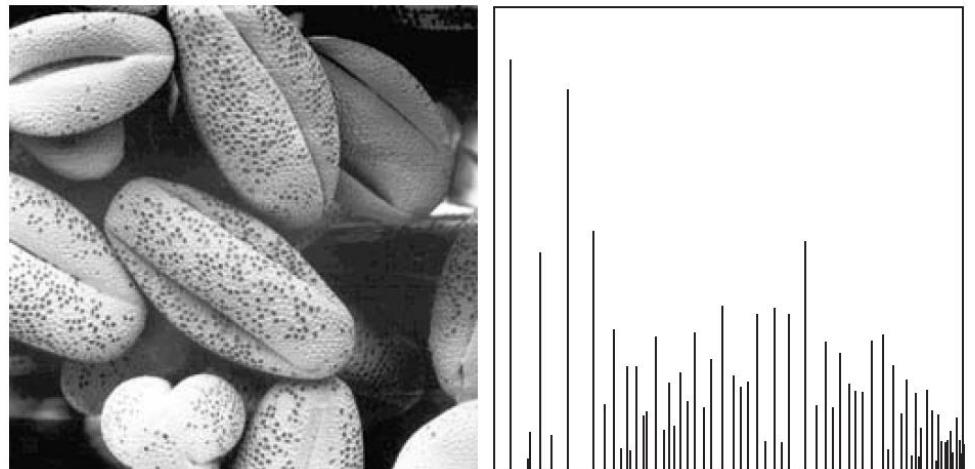
automatic contrast stretching

global transformation function:  
based on intensity distribution of  
entire image

local histogram equalization by  
restricting transformation to pixel  
neighborhoods



after histogram equalization:



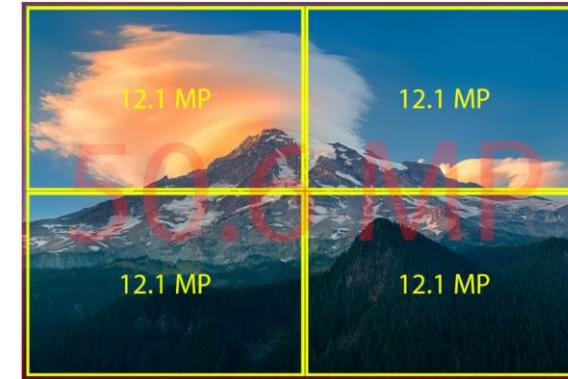
# Spatial Resolution

pixels per unit distance: e.g., pixels per inch (ppi)

→ determines print size of digital image with given pixel count



same resolutions



look at images in screen resolution (property of digital screen)

e.g., iPhone 16 Pro Max:  $2868 \times 1320$  pixels, 6.9 inches diagonal, 460 ppi

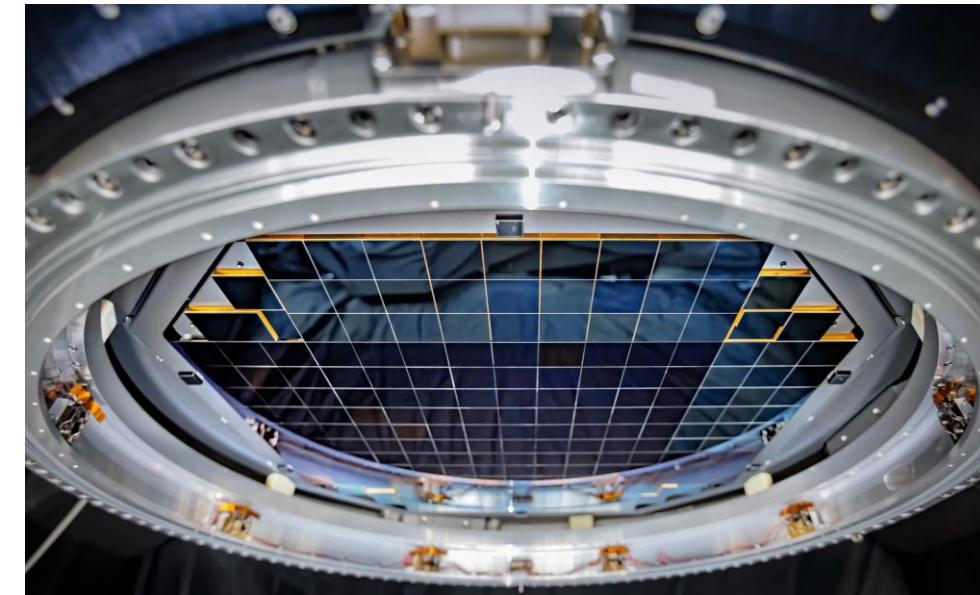
→ either fix number of pixels and calculate figure size, or fix figure size and resample number of pixels (interpolation)

term or printers: dots per inch (dpi) → property of printer

# Example from Astronomy

Vera C. Rubin Observatory

first light expected in 2025



largest digital camera ever made: 3.2-gigapixel CCD



# Intensity Interpolation

# Geometric Transformations

scaling (resizing), rotation, shear, translation, projective warps

approach:

1. relocate individual pixels (without changing its intensity values)
2. calculate new intensity values through interpolation on neighborhood in original image

# Example: Resizing

scaling means changing image size in terms of number of pixels

(Do not confuse this with pixel size, which depends only on the used display or printer.)

- shrinking: sub- or downsampling (pixel deletion)
- zooming: upsampling (pixel replication)

original to be upsampled:



Nearest-neighbor interpolation



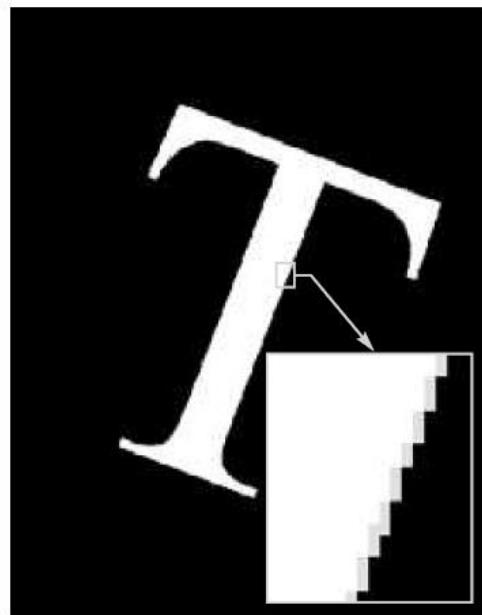
Bilinear interpolation



Bicubic interpolation

# Example: Rotation

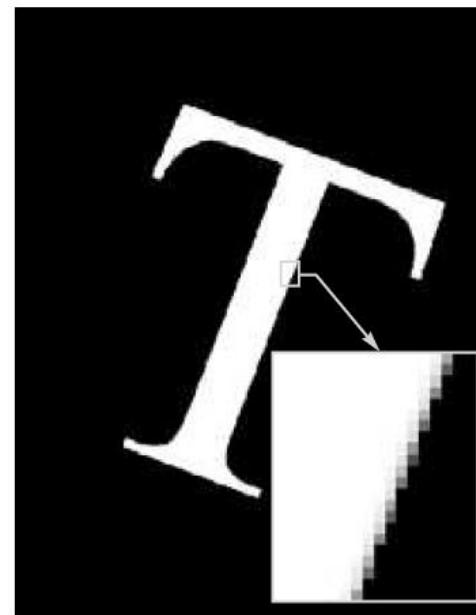
nearest-neighbor  
interpolation:



bilinear interpolation:



bicubic interpolation:



piecewise-polynomial  
function, aka spline

$$v(x, y) = ax + by + cxy + d$$

using 4 neighbors

$$v(x, y) = \sum_{i=0}^3 \sum_{j=0}^3 a_{ij} x^i y^j$$

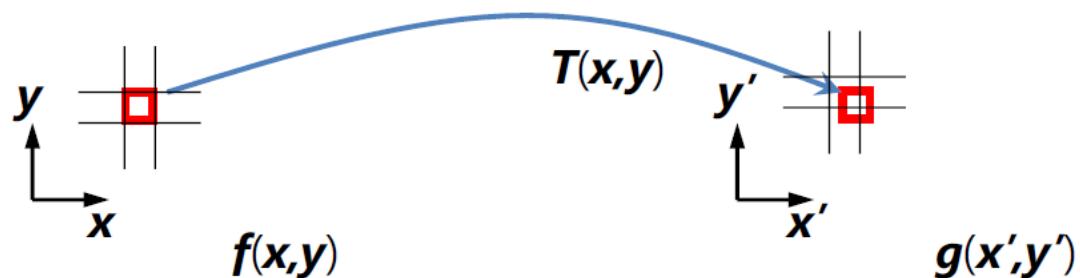
using 16 neighbors

# Mapping

two possibilities:

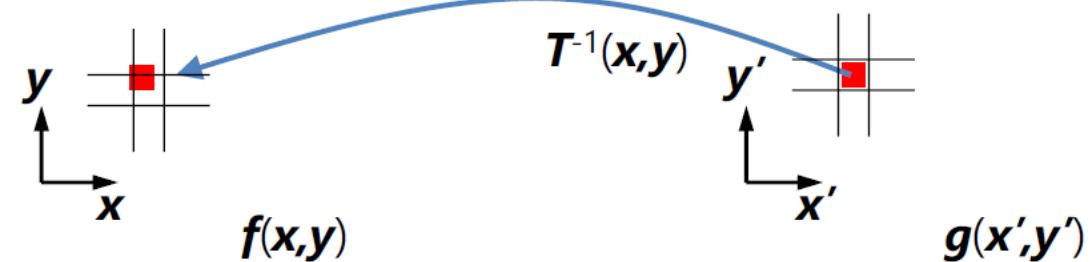
forward mapping:

send each source pixel to its target location



inverse mapping:

get each target pixel from its source location

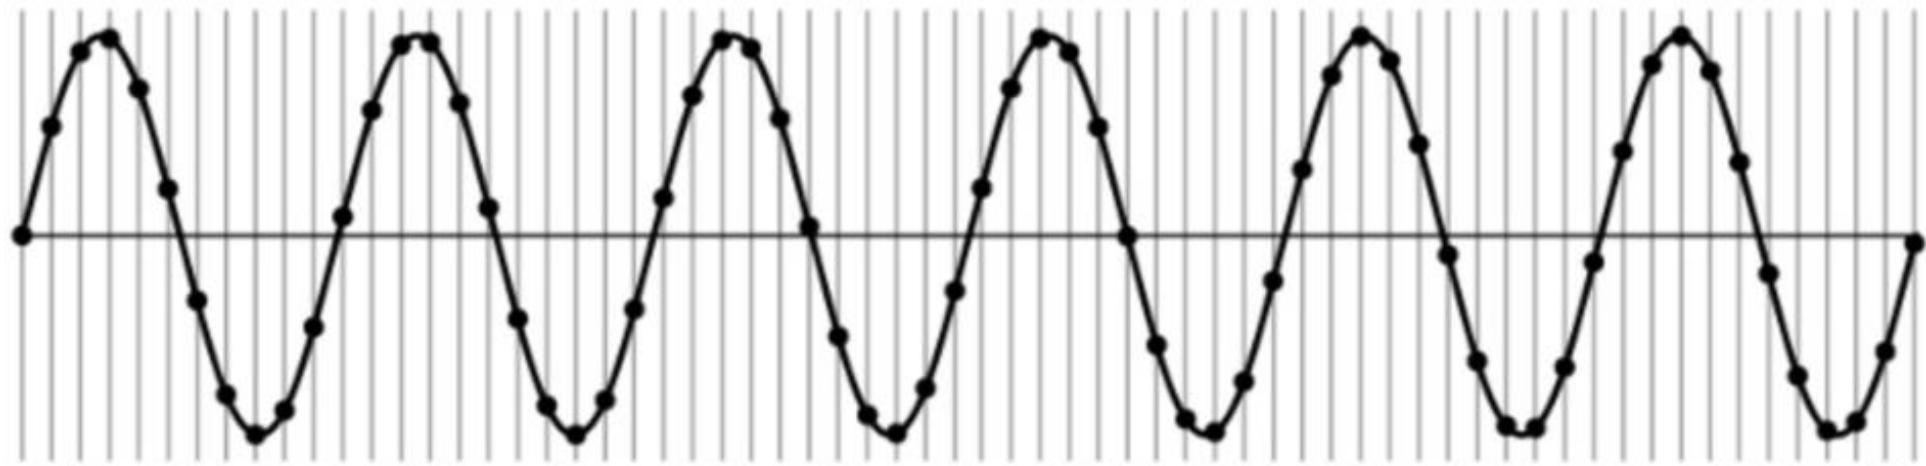


each pixel in original image can contribute  
to more than one pixel in transformed image  
→ can result in holes

each pixel in transformed image can come  
from more than one pixel in original image  
→ resampling from interpolated original image

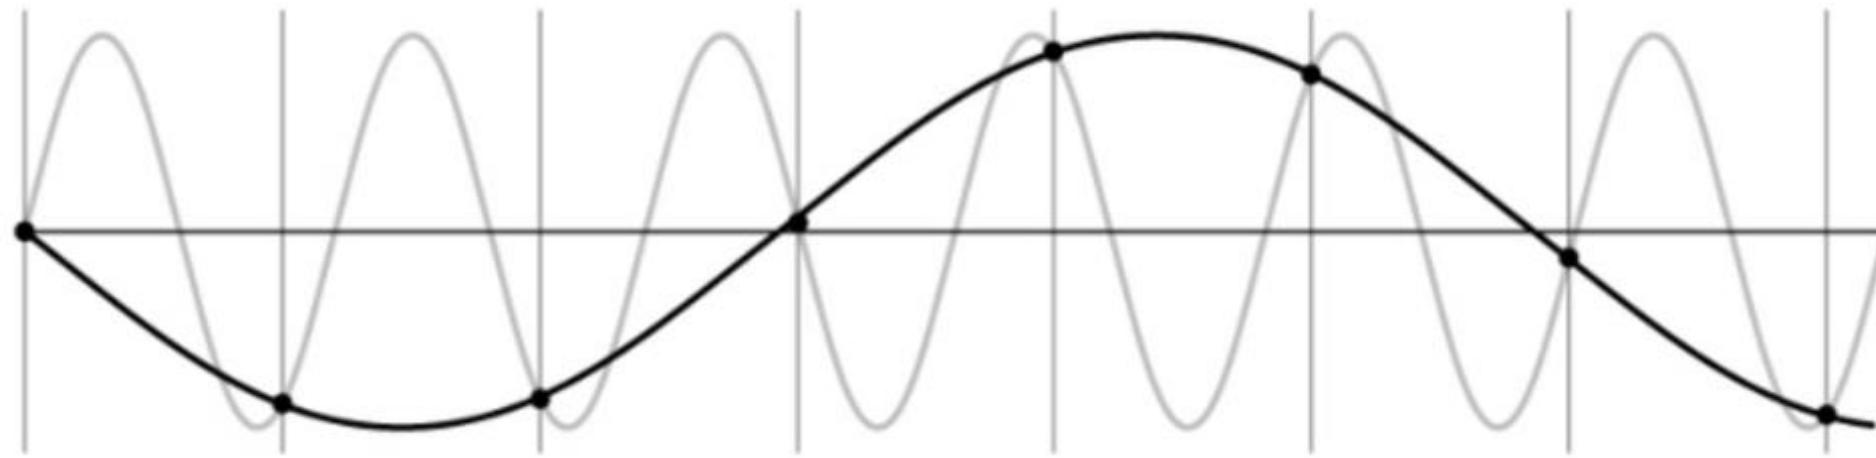
# Aliasing

# Sampling



remember: digital images sample intensity values at each pixel

# Undersampling



**aliasing:** higher-frequency signal disguised as one of lower frequency

# Moiré Effect



aliasing in digital images often manifests itself as ringing patterns

temporal aliasing can be observed as wagon-wheel effect (apparent backward motion) in image sequences (video): frame rate too low with respect to speed of wheel rotation

# Anti-Aliasing in Digital Images

no problem if signal is oversampled:

number of pixels (for both width and height) higher than twice the maximum frequency of image (Nyquist frequency)

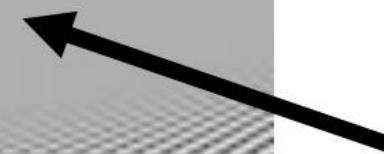
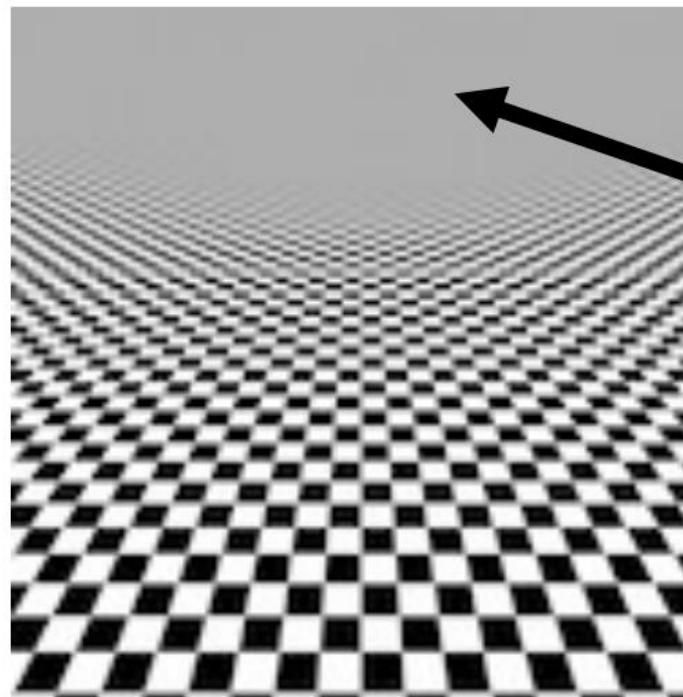
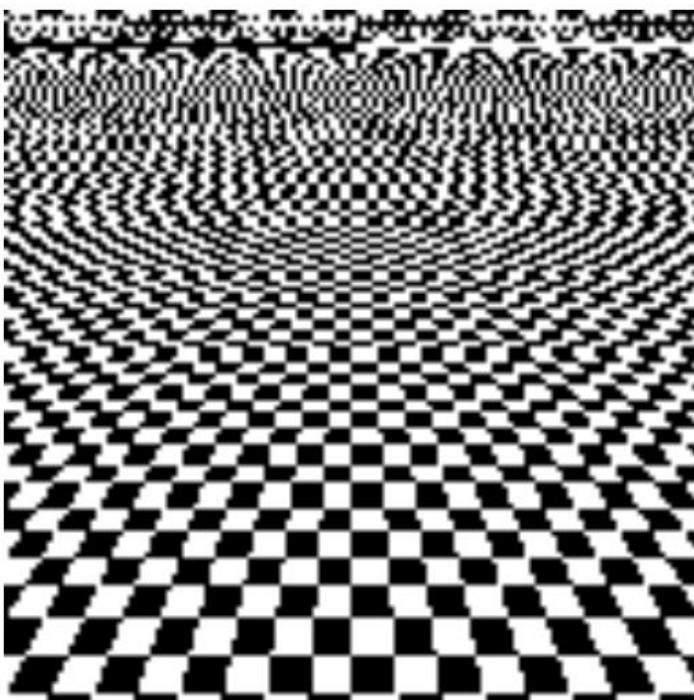
if not oversampled: smooth the signal

→ removing some details (i.e., high frequencies,) that cause aliasing

optical anti-aliasing filter in digital cameras (on top of image sensor): cut out high frequencies (above Nyquist frequency)

# Example: Image Shrinking

smoothing crucial before subsampling (shrinking) to avoid aliasing:  
decreases Nyquist frequency of image (to counter less pixels)



average of black  
and white

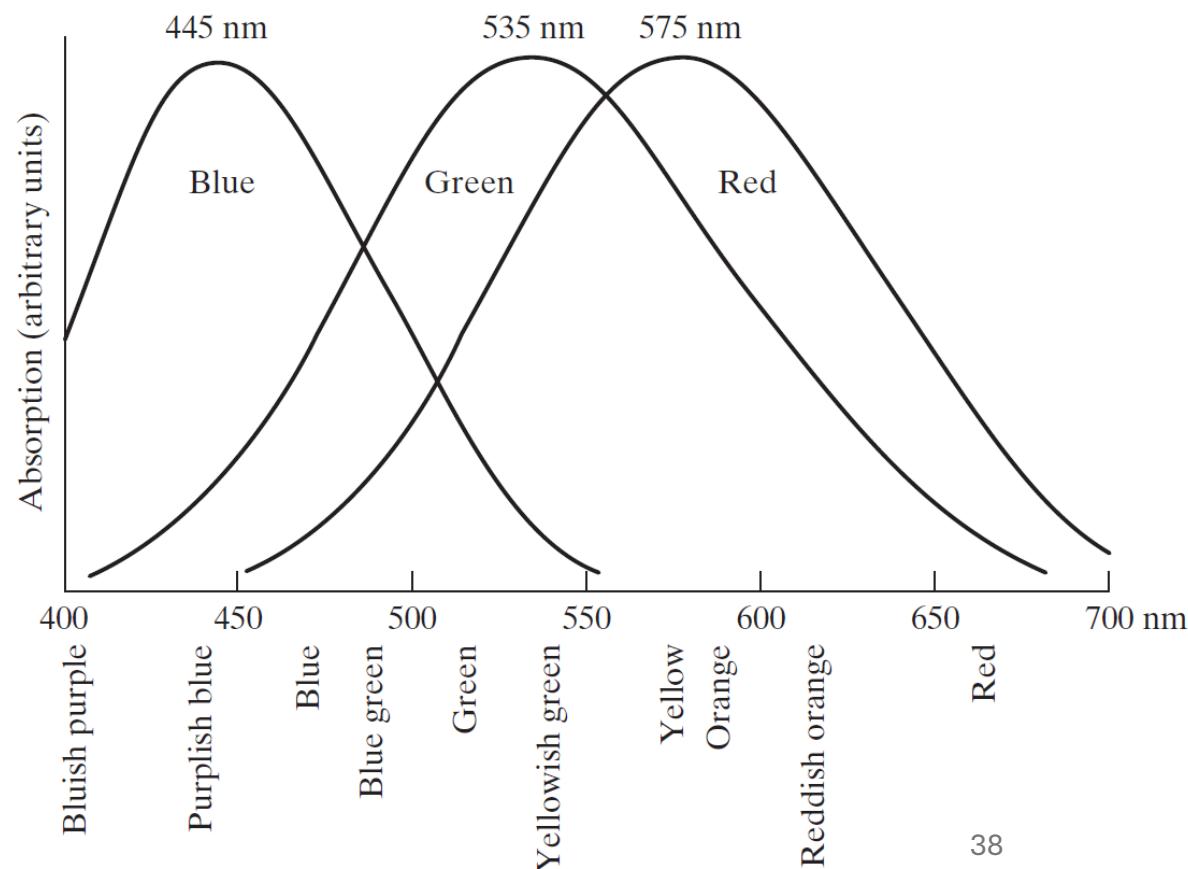
# Color

# Colors in the Human Eye

retina: image sensor of human eye

two kinds of light receptors on surface of retina: rods (monochromatic vision) and cones (color-sensitive cells)

three types of cones: red, green, blue  
→ combinations of red, green, and blue create perceived colors



# Primary Colors

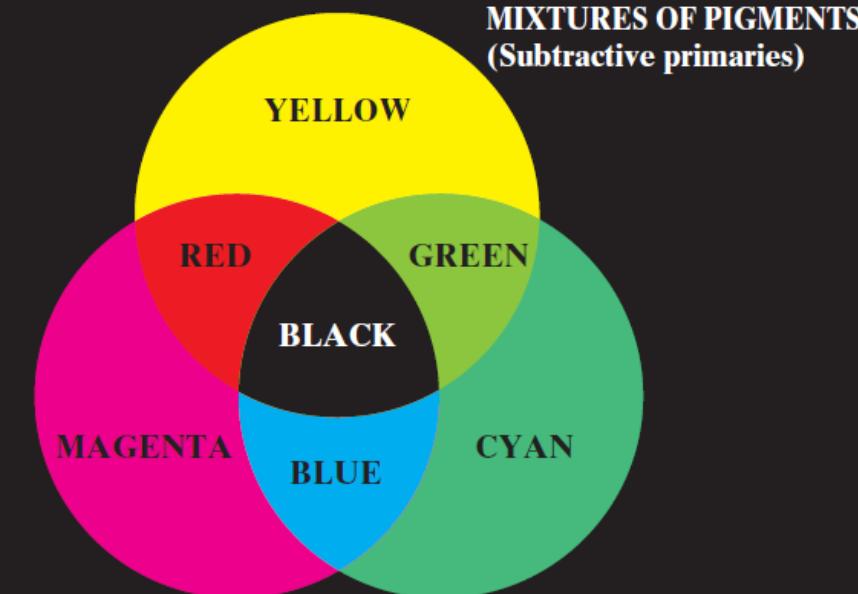
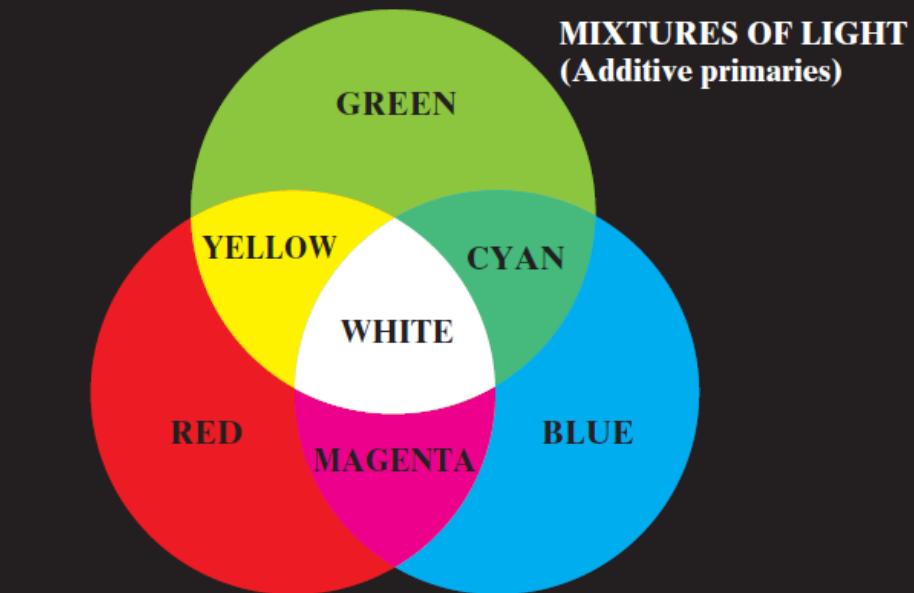
standardization:

use of specific wavelengths for three primary colors red, green, and blue

→ can reproduce (almost) all visible colors by mixing with different intensities

for printing:

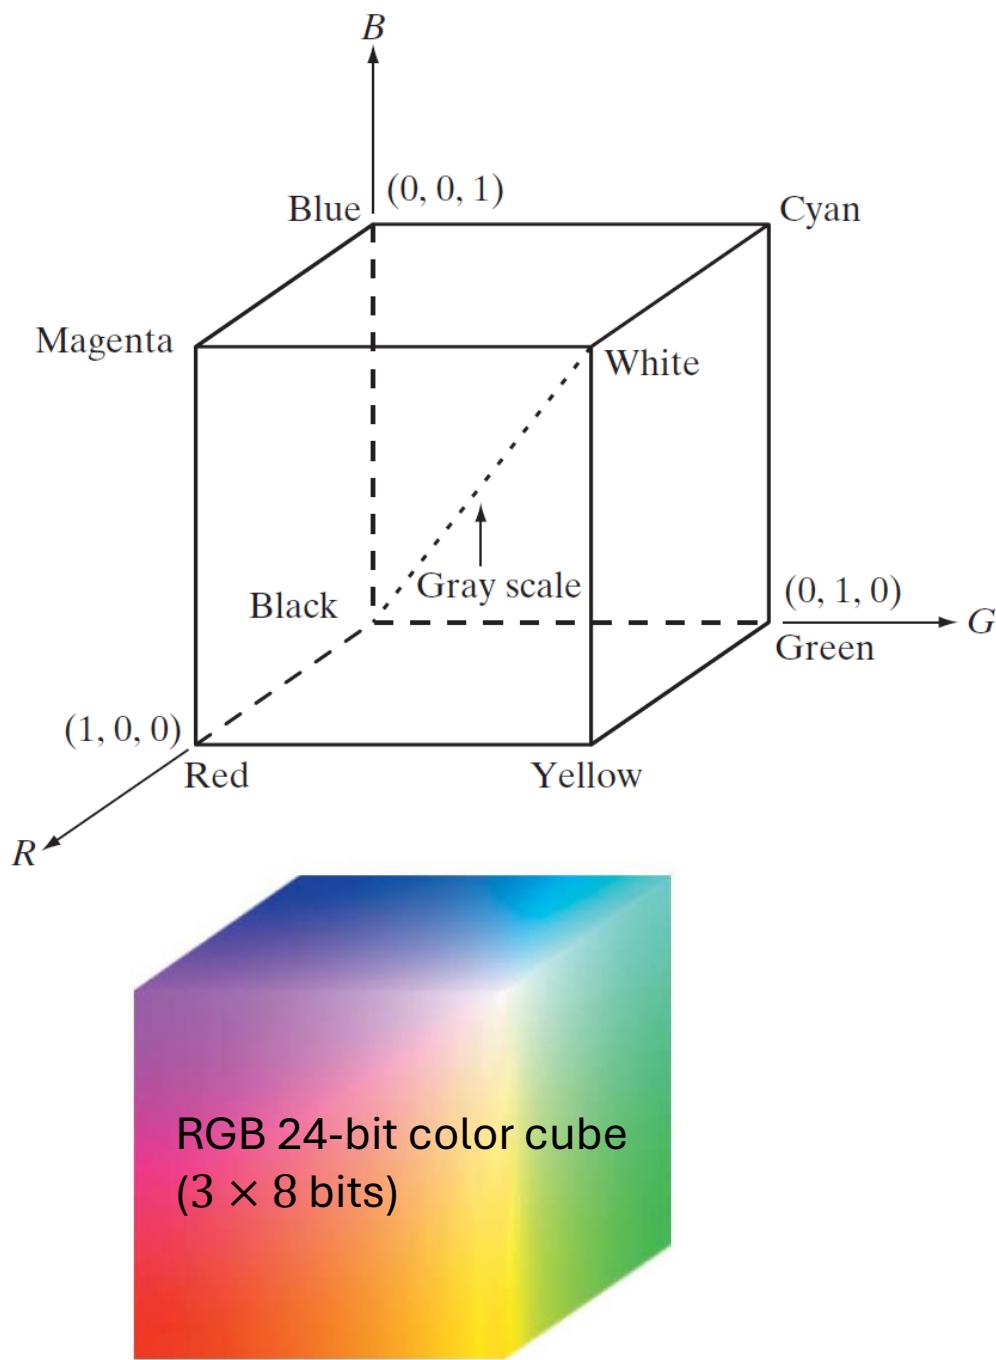
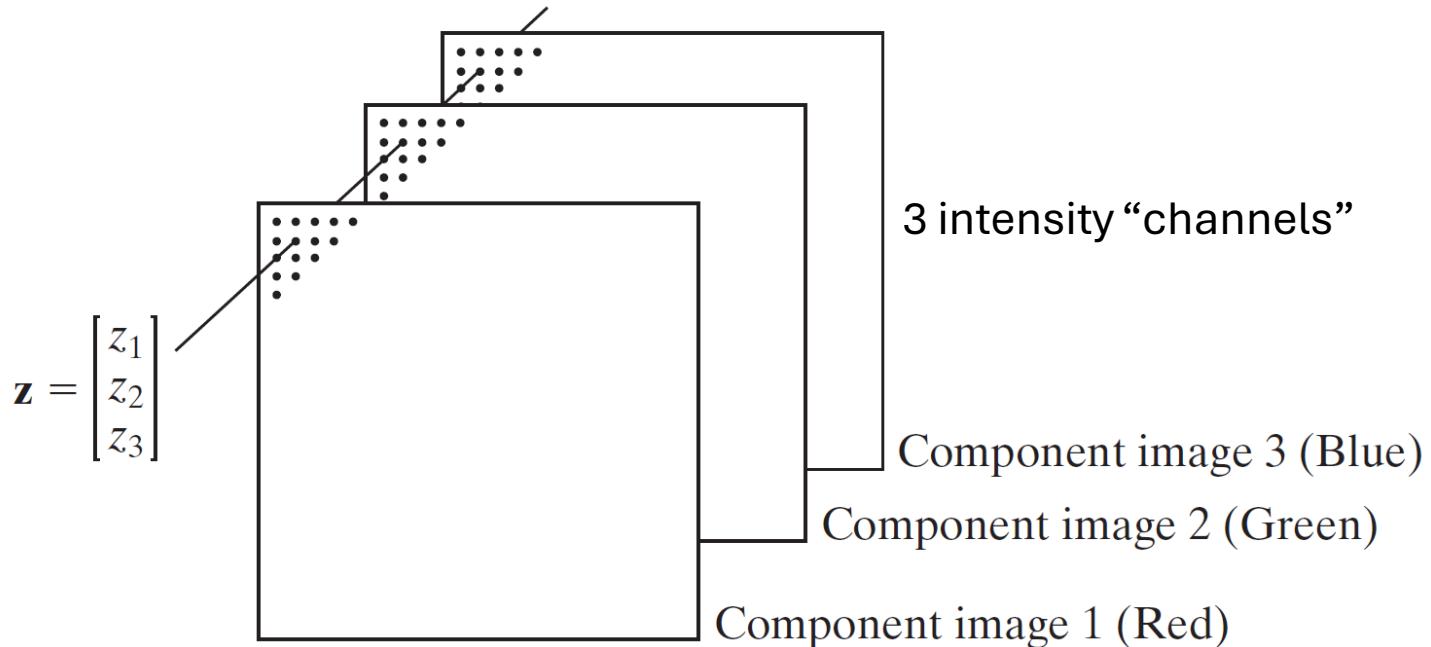
primary colors of pigments each absorb one primary color of light



# RGB Color Model

color model for monitors and cameras  
(other color models: CMYK, HSV, ...)

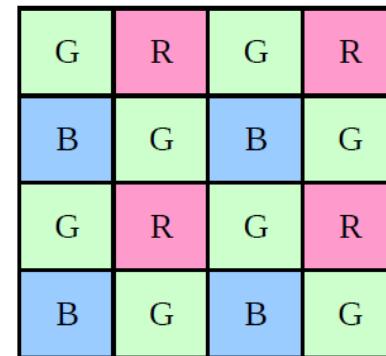
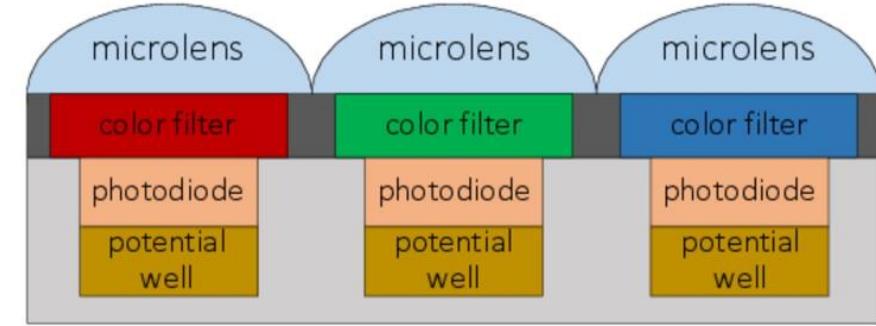
color pixels:



# Color Sensing

single image sensor with one of three color filters in front of each pixel

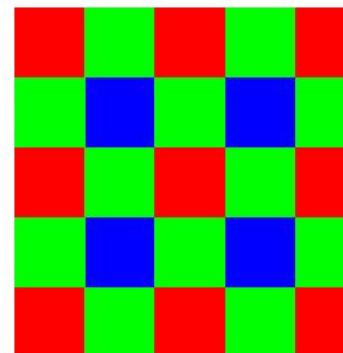
estimation of missing color values for each pixel by interpolation from neighboring pixels  
→ full-color image



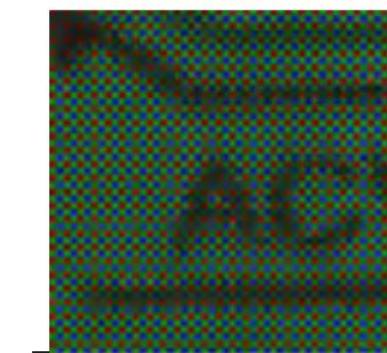
Bayer RGB Pattern

rGb	Rgb	rGb	Rgb
rgB	rGb	rgB	rGb
rGb	Rgb	rGb	Rgb
rgB	rGb	rgB	rGb

Interpolated Pixels



Bayer Pattern  
(Color Filter Mosaic)



Raw Image



Interpolated Image

# Neighborhood Operators

## Single-Pixel/Point Operation



point processing  
intensity transformations

## Neighborhood Operation



“filtering”

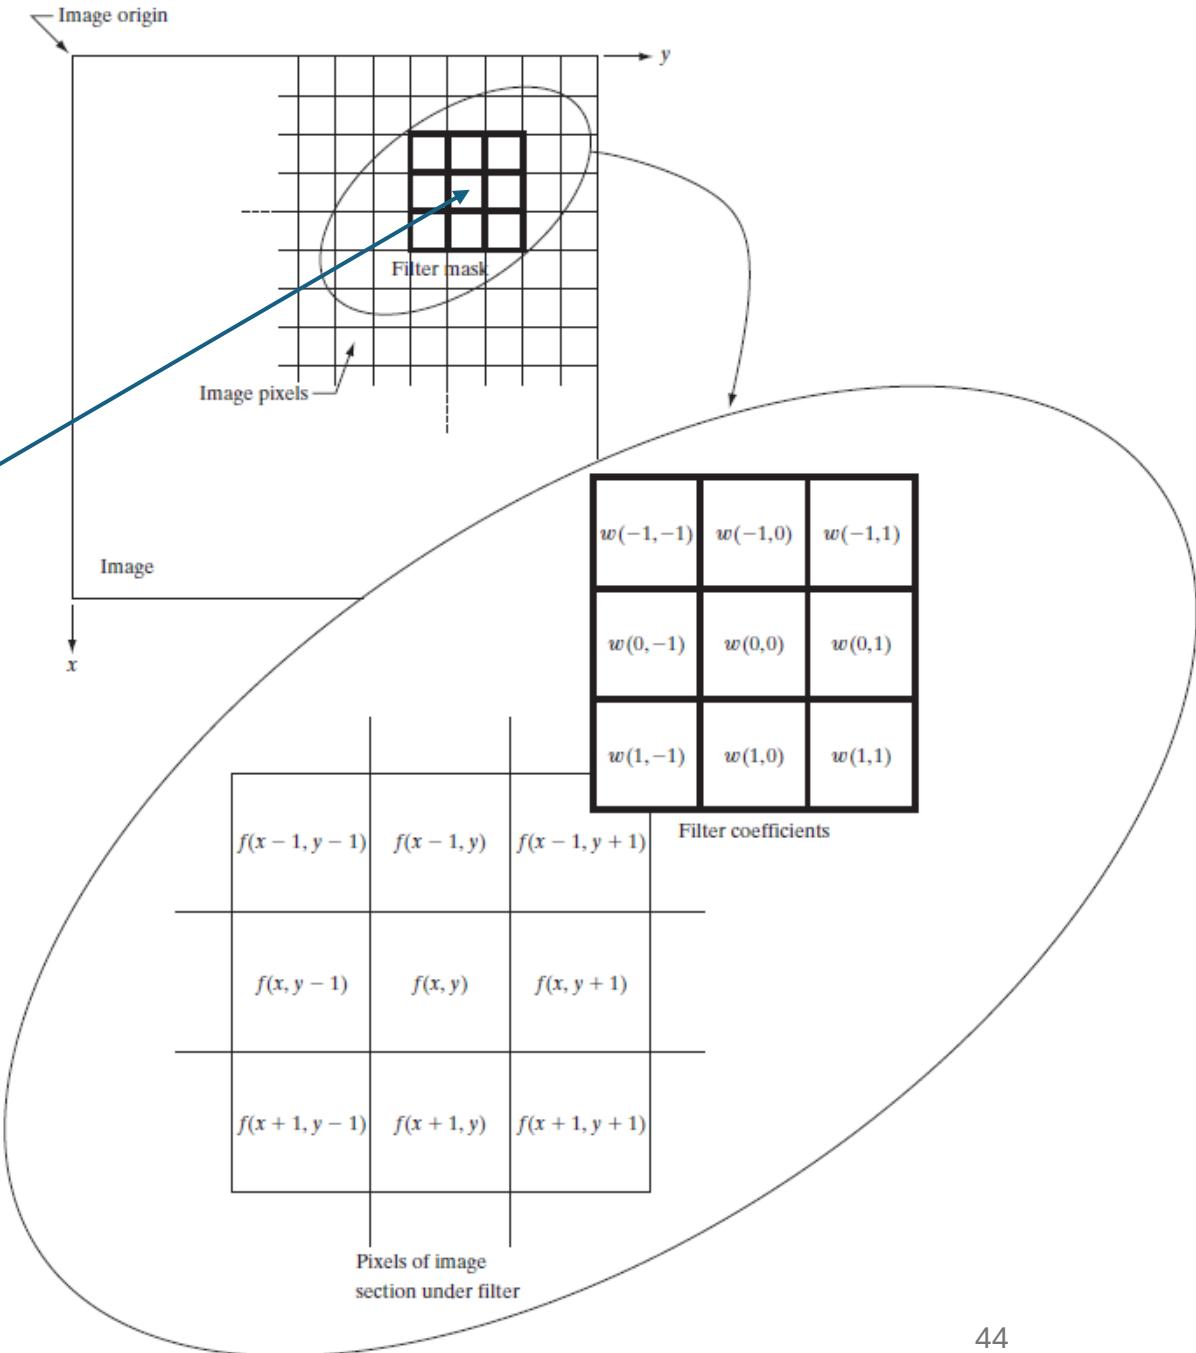
# Spatial Filtering

linear shift-invariant operators:

- replace each pixel's intensity by weighted sum of its neighborhood
- weights determined by a kernel (aka mask or filter)
- same kernel **shifted** to all pixels

for example, filter of size  $3 \times 3$ :

$$g(x, y) = \sum_{s=-1}^1 \sum_{t=-1}^1 w(s, t)f(x + s, y + t)$$



We set the filter coefficients manually here.  
Later, we will use a CNN for this.

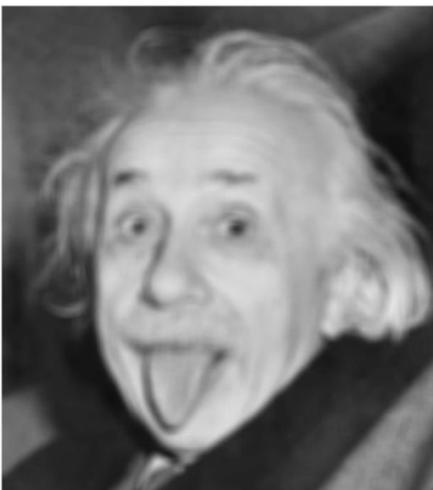
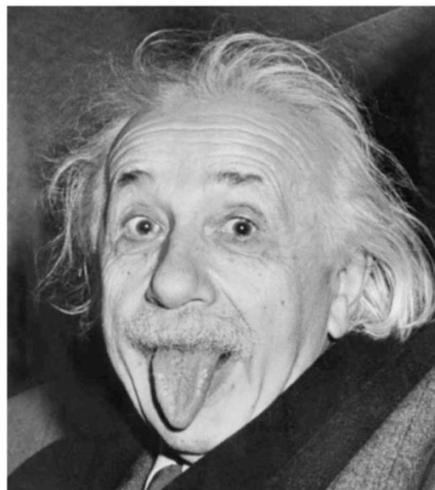
# Example: Box Filter

$$w = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$m \times m: w = \frac{1}{m^2} \begin{bmatrix} 1 & \dots & 1 \\ \vdots & \ddots & \vdots \\ 1 & \dots & 1 \end{bmatrix}$$

the larger the kernel,  
the heavier the  
smoothing effect

results in local (moving) averages  
→ blurring (smoothing) effect:



- without normalization resulting image would be saturated
- the larger the filter size, the heavier the blurring effect
- filters of even size possible as well, but odd-sized ones easier to deal with

# Definition: Correlation vs Convolution

using our example with filter of size  $3 \times 3$

correlation: 
$$g(x, y) = \sum_{s=-1}^1 \sum_{t=-1}^1 w(s, t)f(x + s, y + t)$$

convolution: 
$$g(x, y) = \sum_{s=-1}^1 \sum_{t=-1}^1 w(s, t)f(x - s, y - t)$$

corresponds to flip of filter

- convolution of a kernel with a discrete unit impulse yields copy of the kernel at location of the impulse (and impulse filter yields image copy)
- corresponding correlation instead yields rotated version of the kernel (or rotated image)





# Padding

On last slides, we did not let the filter extend beyond image borders.  
→ boundary effects: resulting image smaller than original

zero-padding to avoid this:

$$\begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{matrix} * \begin{matrix} 1 & 2 & -1 \\ 1 & 1 & -1 \\ -2 & -1 & 0 \end{matrix} = \begin{matrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & -1 & 0 \\ 0 & 1 & 1 & -1 & 0 \\ 0 & -2 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{matrix}$$

zero-padding darkens  
the image edges

- other padding options:
- set to constant border value
  - repeat edge values
  - ...

can also be more zero rows/columns (depending on image and filter size)

# Other Linear Smoothing Filters

when constructing filters (not only smoothing filters, but in general), look for separable ones:

allow to compute 2D convolution as two subsequent 1D convolutions (horizontal and vertical)

for image of size  $M \times M$  and kernel of size  $N \times N$ :

- $M^2 \cdot N^2$  operations for non-separable filter
  - $M^2 \cdot 2 \cdot N$  operations for separable filter
- significant speedup possible

sampled from 2D Gaussian function  
(larger  $\sigma$  requires larger kernel)

bilinear filter  
(tent):

$\frac{1}{16}$	1	2	1
	2	4	2
	1	2	1

$$\frac{1}{256}$$

Gaussian filter:

1	4	6	4	1
4	16	24	16	4
6	24	36	24	6
4	16	24	16	4
1	4	6	4	1

$\frac{1}{4}$	1	2	1
	1	4	1

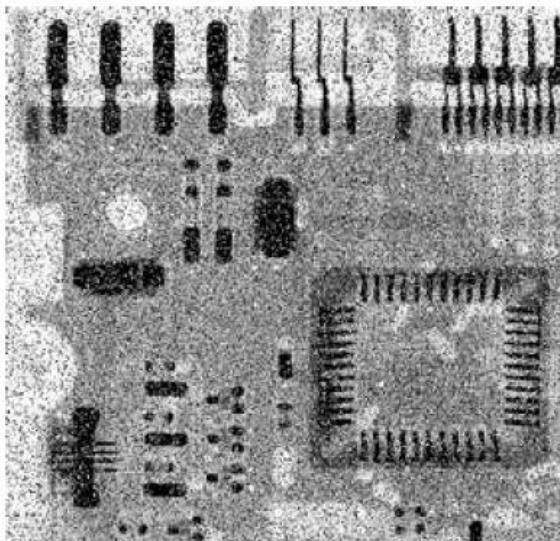
$\frac{1}{16}$	1	4	6	4	1
	1	4	6	4	1

corresponding horizontal 1D kernels

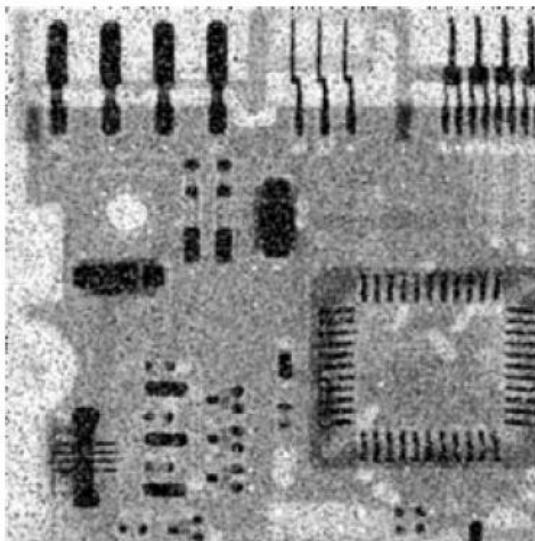
# Median Filter

- same kind of neighborhood-shifting over image
- but instead of weighted average, use median of intensity values in pixel neighborhood (no filter coefficients) → non-linear filter
- very effective to reduce outliers, e.g., shot or salt-and-pepper noise

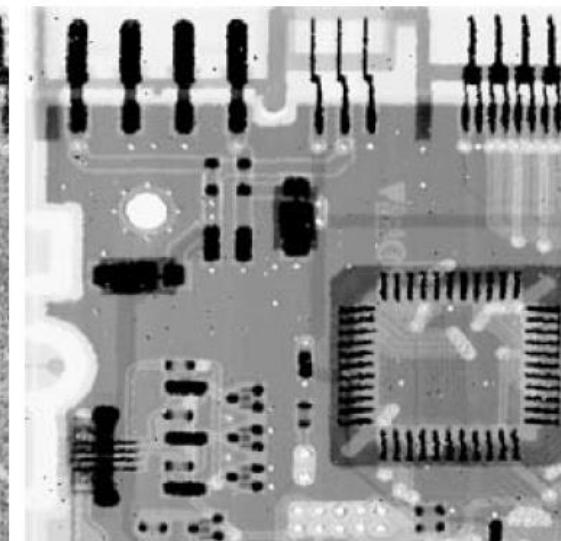
original:



average-filtered ( $3 \times 3$ ):



median-filtered ( $3 \times 3$ ):

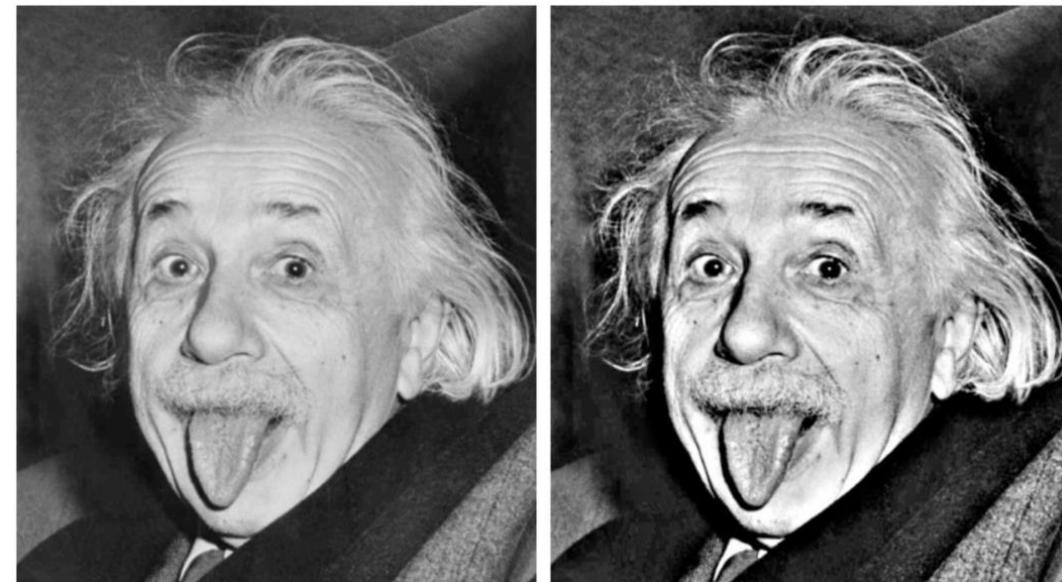


# Sharpening Filters

goal of smoothing filters:  
removal of small details, noise reduction

goal of sharpening filters:  
highlight intensity transitions (e.g., edges)

averaging analogous to integration  
→ sharpening by differentiation



# First- and Second-Order Derivatives

partial derivatives in discrete form:

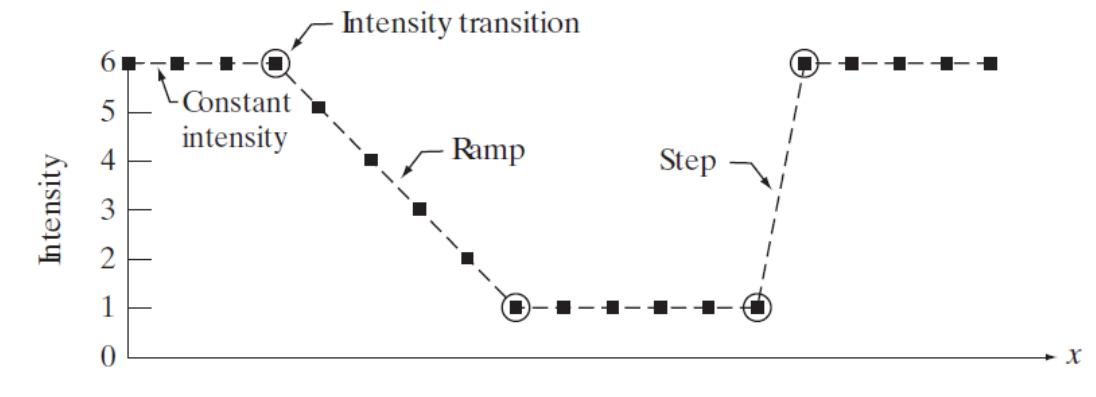
$$\frac{\partial f}{\partial x} = f(x + 1, y) - f(x, y)$$

$$\frac{\partial f}{\partial y} = f(x, y + 1) - f(x, y)$$

$$\frac{\partial^2 f}{\partial x^2} = f(x + 1, y) + f(x - 1, y) - 2f(x, y)$$

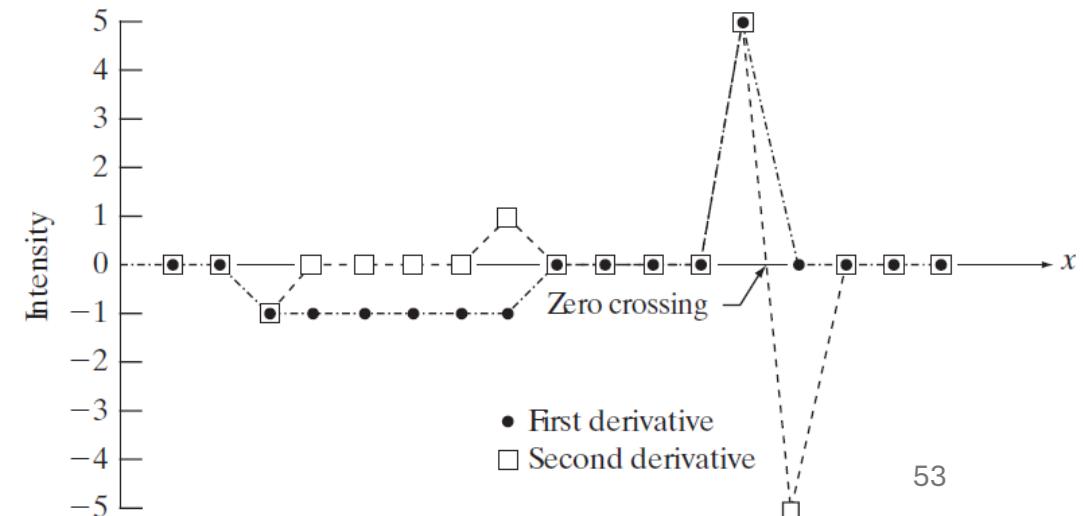
$$\frac{\partial^2 f}{\partial y^2} = f(x, y + 1) + f(x, y - 1) - 2f(x, y)$$

1D visualization:



	6	6	6	6	5	4	3	2	1	1	1	1	1	1	6	6	6	6	6
Scan line	6	6	6	6	5	4	3	2	1	1	1	1	1	1	6	6	6	6	6
1st derivative	0	0	-1	-1	-1	-1	-1	0	0	0	0	0	0	0	5	0	0	0	0

	0	0	-1	0	0	0	0	1	0	0	0	0	0	0	5	0	0	0	0
2nd derivative	0	0	-1	0	0	0	0	0	1	0	0	0	0	0	5	-5	0	0	0



# Approximated Image Gradient

alternative formulation (1D):  $\frac{\partial f}{\partial x} = \frac{f(x+1) - f(x-1)}{2}$

→ 1D derivative filter:

$$\begin{array}{|c|c|c|}\hline 1 & 0 & -1 \\ \hline\end{array}$$

Sobel filters:

smoothing

$$\begin{array}{|c|}\hline 1 \\ \hline 2 \\ \hline 1 \\ \hline\end{array}$$

horizontal

$$\begin{array}{|c|c|c|}\hline -1 & 0 & 1 \\ \hline\end{array}$$

$$\begin{array}{|c|c|c|}\hline 1 & 0 & -1 \\ \hline 2 & 0 & -2 \\ \hline 1 & 0 & -1 \\ \hline\end{array}$$

$$\cdot \frac{1}{8}$$

(convolution kernels)

construct

separable filters:

vertical

$$\begin{array}{|c|}\hline -1 \\ \hline 0 \\ \hline 1 \\ \hline\end{array}$$

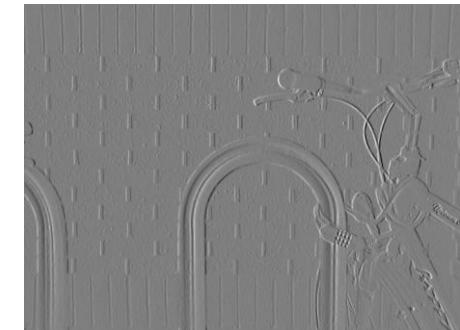
smoothing

$$\begin{array}{|c|c|c|}\hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline\end{array}$$

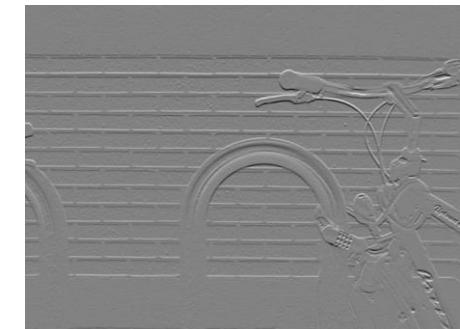
$$\cdot \frac{1}{8}$$

→ gradient with smoothing

convolved with an image:



$$\frac{\partial f}{\partial x}$$



$$\frac{\partial f}{\partial y}$$

# Magnitude of Image Gradient

magnitude/amplitude of gradient:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

magnitude is isotropic

for sharpening:  
add this to original



# Different Gradient Filters

Gradient	Roberts	Prewitt	Sobel (3x3)	Sobel (5x5)
$\frac{\partial I}{\partial x}$	$\begin{matrix} 0 & 1 \\ -1 & 0 \end{matrix}$	$\begin{matrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{matrix}$	$\begin{matrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{matrix}$	$\begin{matrix} -1 & -2 & 0 & 2 & 1 \\ -2 & -3 & 0 & 3 & 2 \\ -3 & -5 & 0 & 5 & 3 \\ -2 & -3 & 0 & 3 & 2 \\ -1 & -2 & 0 & 2 & 1 \end{matrix}$
$\frac{\partial I}{\partial y}$	$\begin{matrix} 1 & 0 \\ 0 & -1 \end{matrix}$	$\begin{matrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{matrix}$	$\begin{matrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{matrix}$	$\begin{matrix} 1 & 2 & 3 & 2 & 1 \\ 2 & 3 & 5 & 3 & 2 \\ 0 & 0 & 0 & 0 & 0 \\ -2 & -3 & -5 & -3 & -2 \\ -1 & -2 & -3 & -2 & -1 \end{matrix}$

Good Localization  
Noise Sensitive  
Poor Detection

Poor Localization  
Less Noise Sensitive  
Good Detection

the larger ones add  
smoothing components  
(Sobel: convolution  
with bilinear filter)

# Alternative: Laplacian

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

filter mask:

0	1	0
1	-4	1
0	1	0

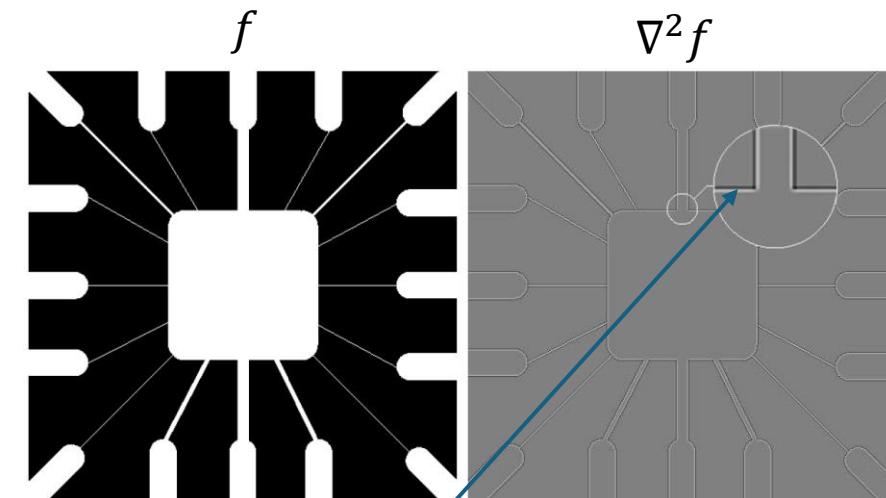
$$= f(x+1, y) + f(x-1, y) + f(x, y+1) + f(x, y-1) - 4f(x, y)$$

isotropic for rotations in increments of  $90^\circ$

to yield isotropic results for rotations in increments of  $45^\circ$ :

again, add to (or rather subtract from when using this filter definition here) original to get sharpened image:

$$g(x, y) = f(x, y) - \nabla^2 f(x, y)$$



scaled to visualize  
negative values

zero-crossings  $\rightarrow$  edges

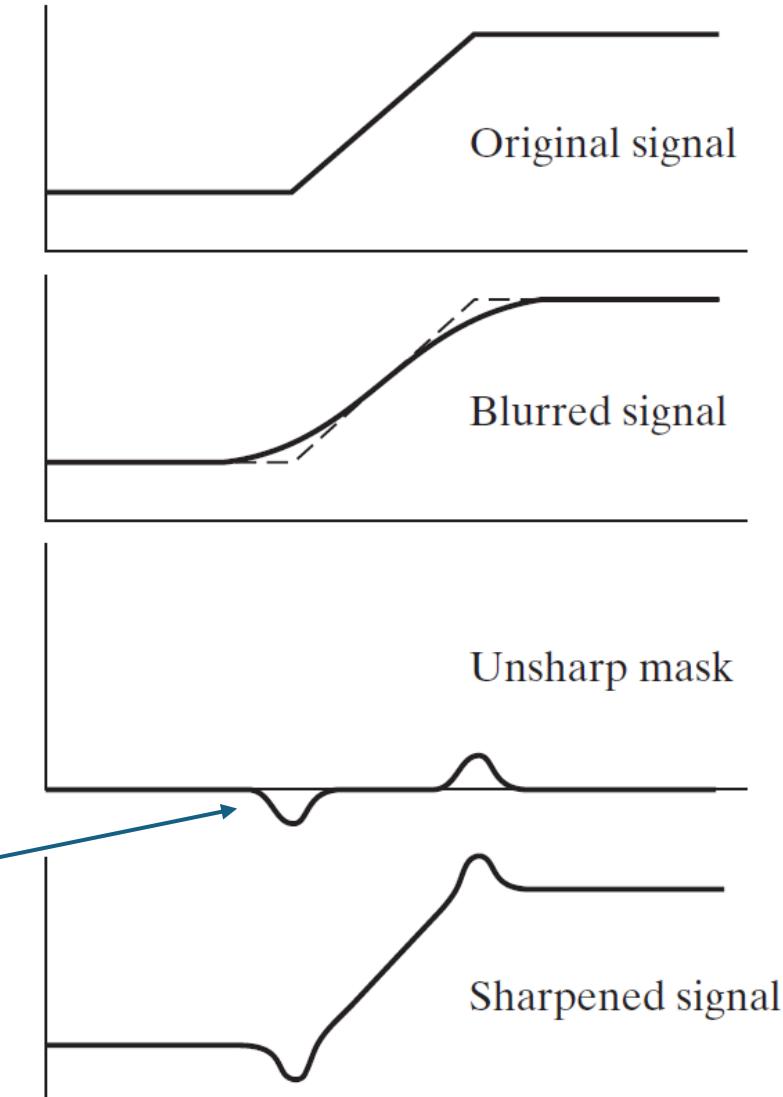
1	1	1
1	-8	1
1	1	1

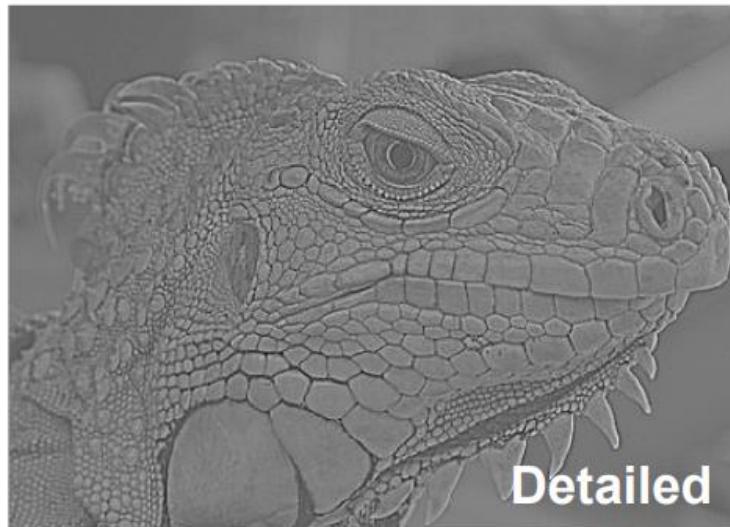
# Another Alternative: Unsharp Masking

use of smoothing filters for sharpening

1. smoothing of original image (e.g., by Gaussian filter)
2. subtract the blurred image from the original → unsharp mask (similar to second-order derivative)
3. add unsharp mask to original image → sharpened image

similar to second-order derivative



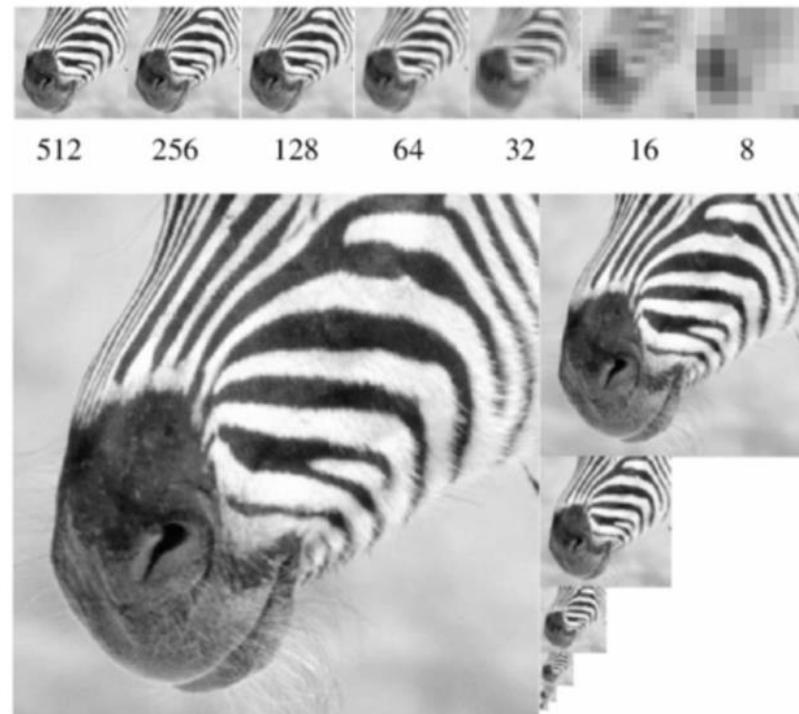
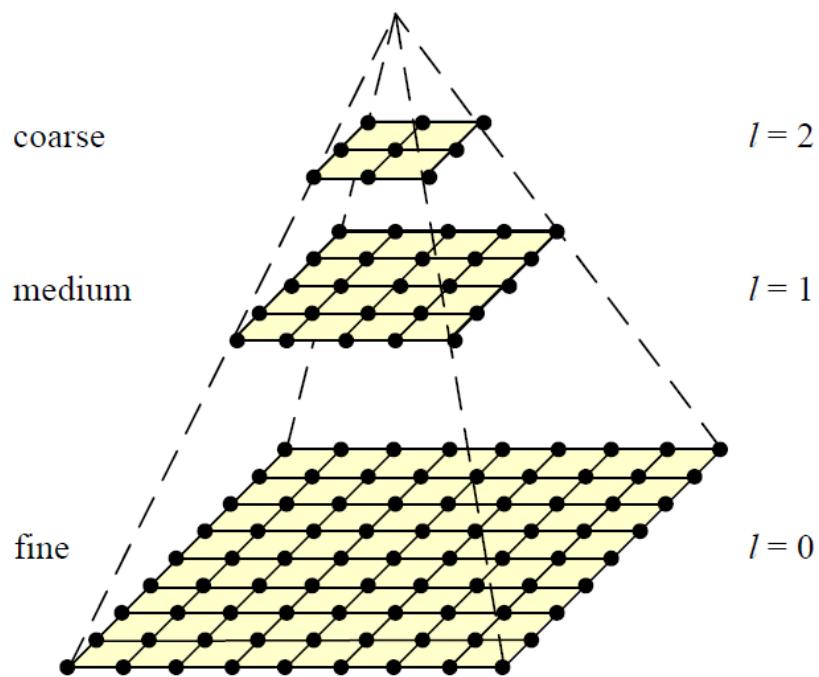


# Image Pyramids

# Gaussian Image Pyramid

sequence of sub-sampled images (one use case: features at different scales)

apply smoothing before each down-sample step (to avoid aliasing), typically by convolution with binomial kernel (converging to Gaussian)

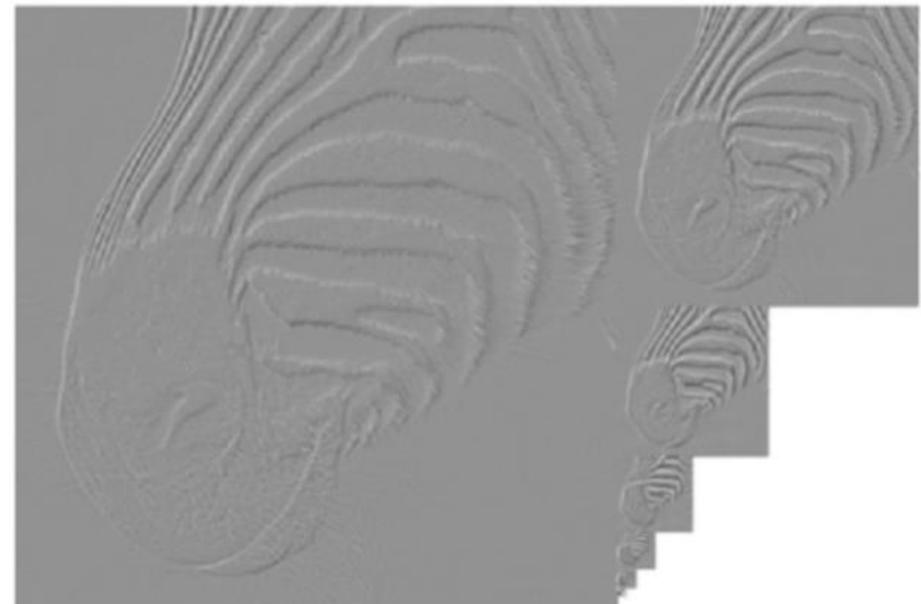


# Laplacian Image Pyramid

retain only residuals at different levels (allows later upsampling, starting from highest-level image)



512 256 128 64 32 16 8



difference of Gaussians approximate Laplacian:

